

# Convolutional Neural Networks II

CS 4391 Introduction Computer Vision

Instructor Yu Xiang

The University of Texas at Dallas

Some slides of this lecture are courtesy Stanford CS231n

# Fully Connected Layer

- What is the drawback of only using fully connected layers?

$$y = Wx$$

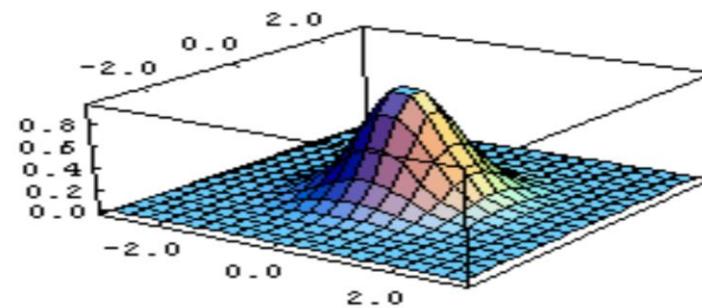
- Consider an image with  $640 \times 480$ 
  - $x$  is with dimension 307,200
  - The weight matrix of the fully connect layer is too large

# Convolutional Layers

- Consist of convolutional filters
- Share weights among different image locations

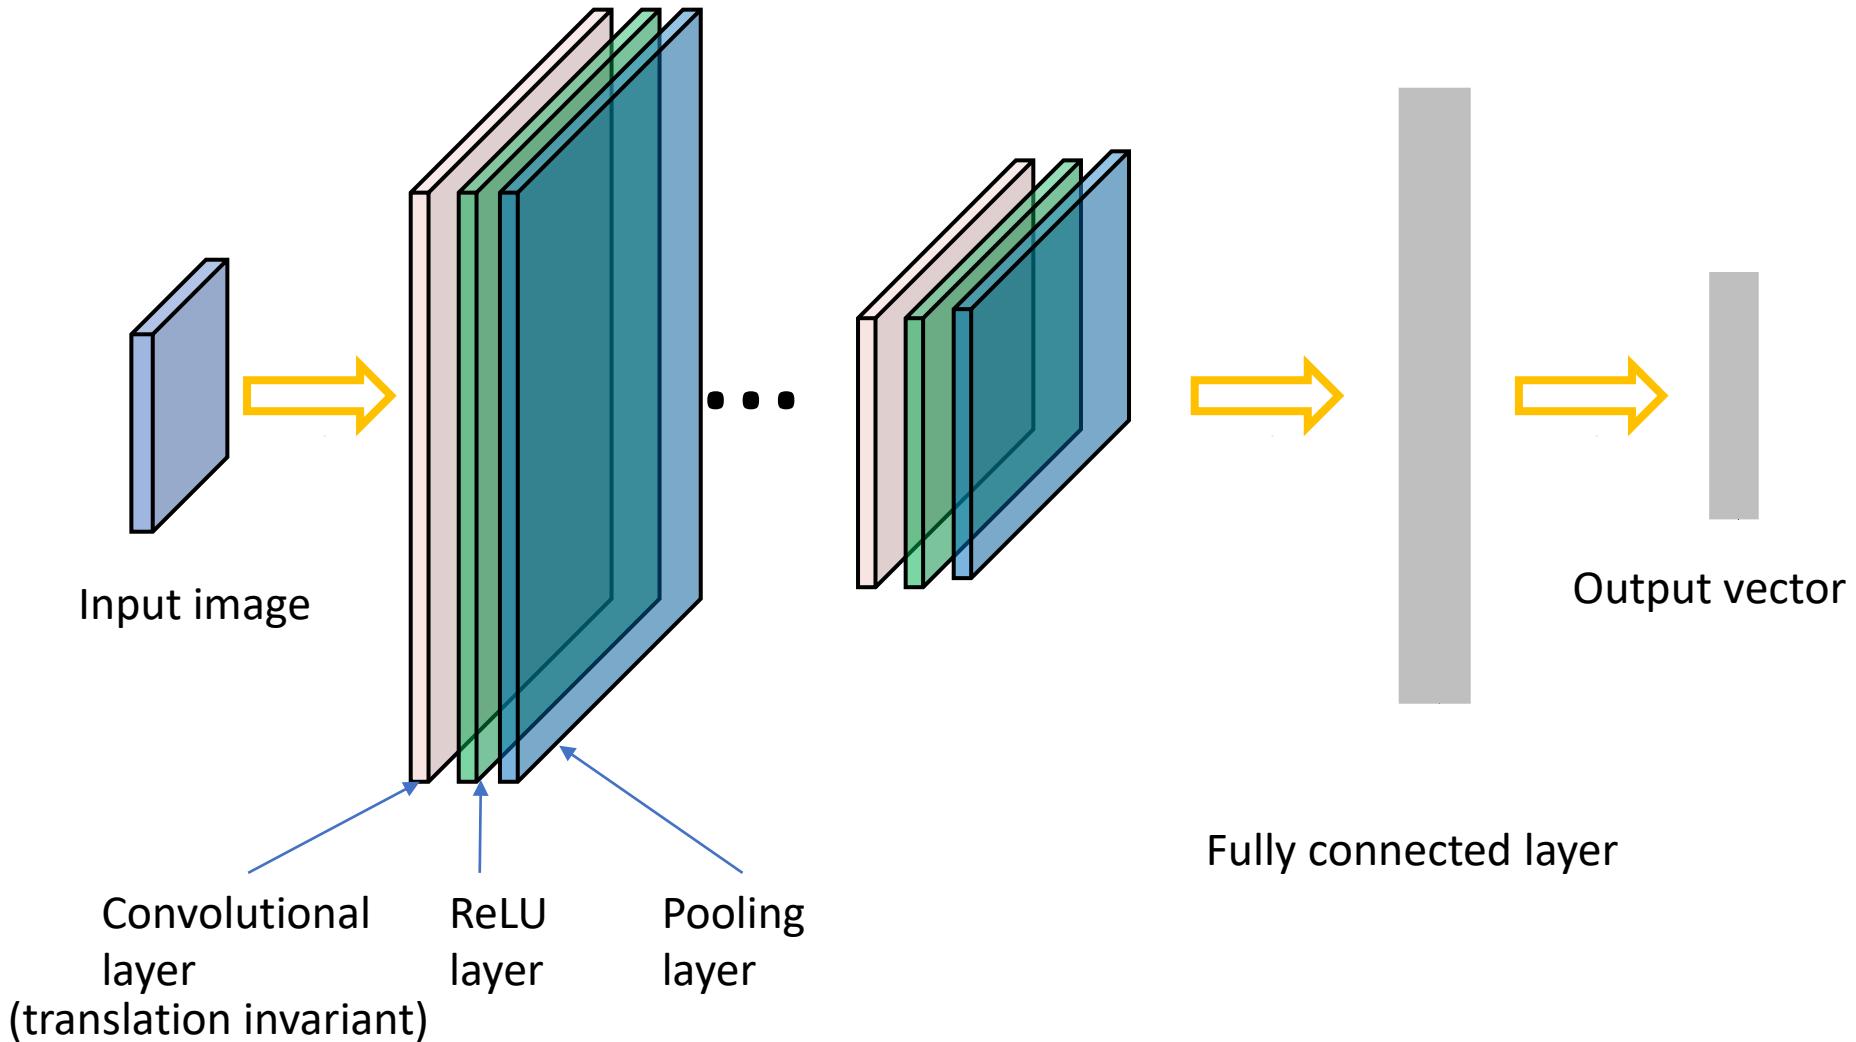
$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Gaussian Filter

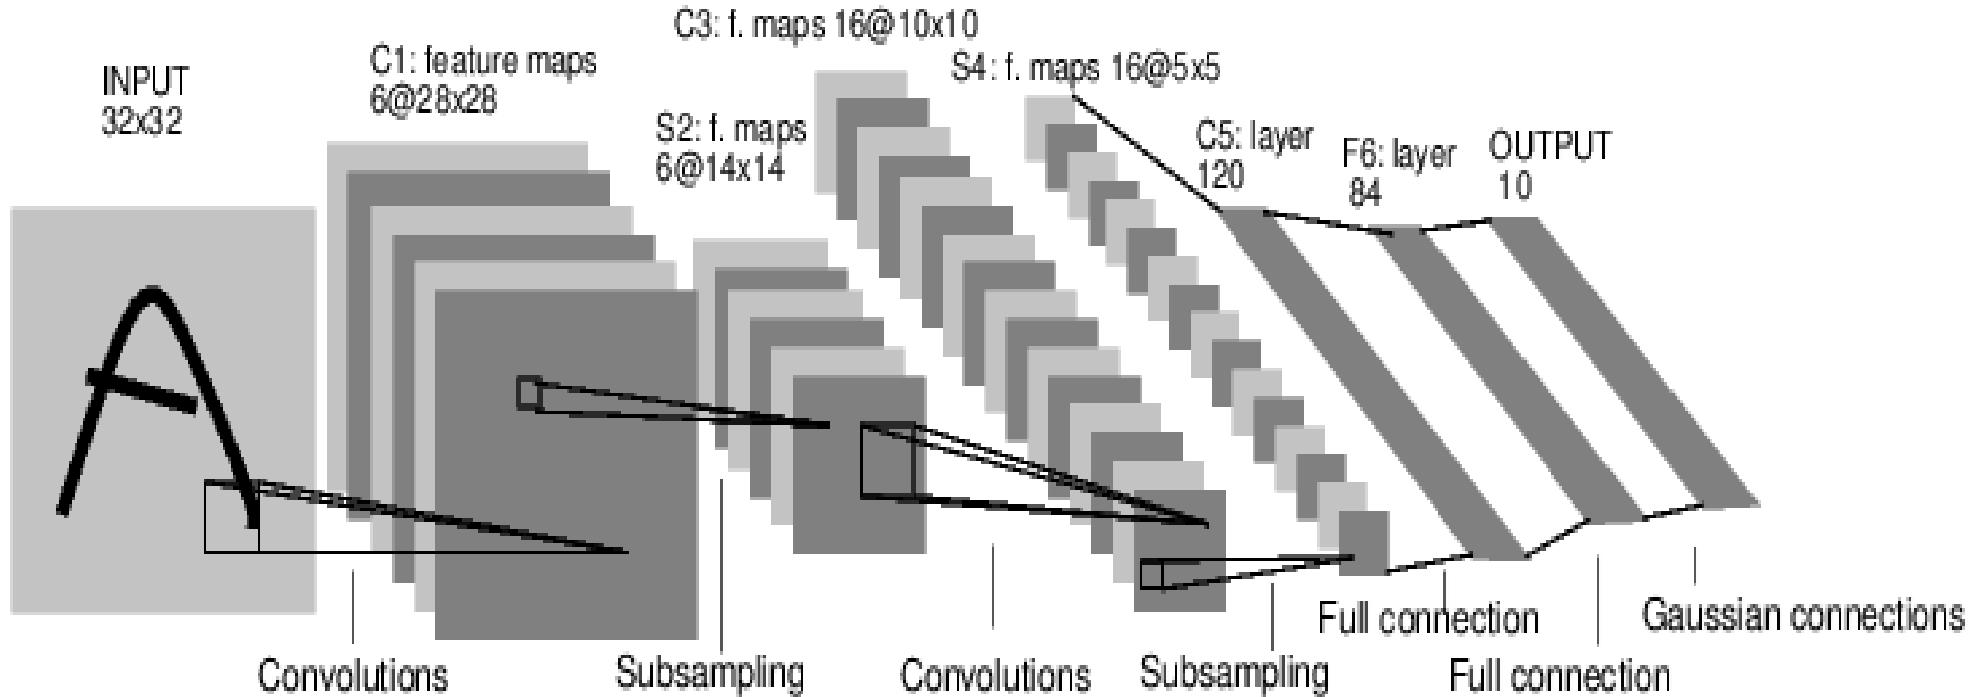


Learn the weights!

# Convolutional Neural Networks



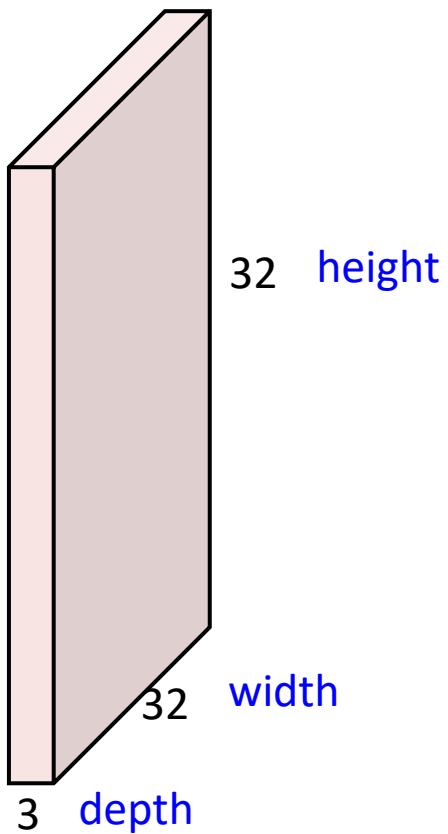
# Convolutional Neural Networks



[LeNet-5, LeCun 1980]

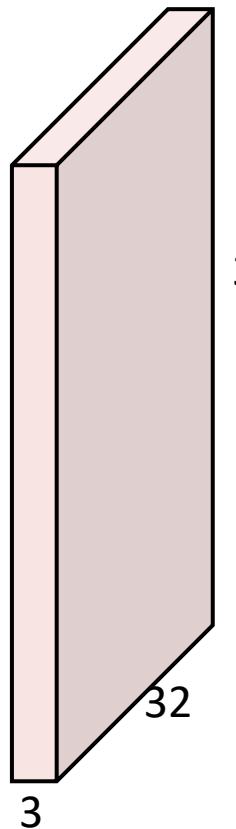
# Convolutional Layer

32x32x3 image



# Convolutional Layer

32x32x3 image

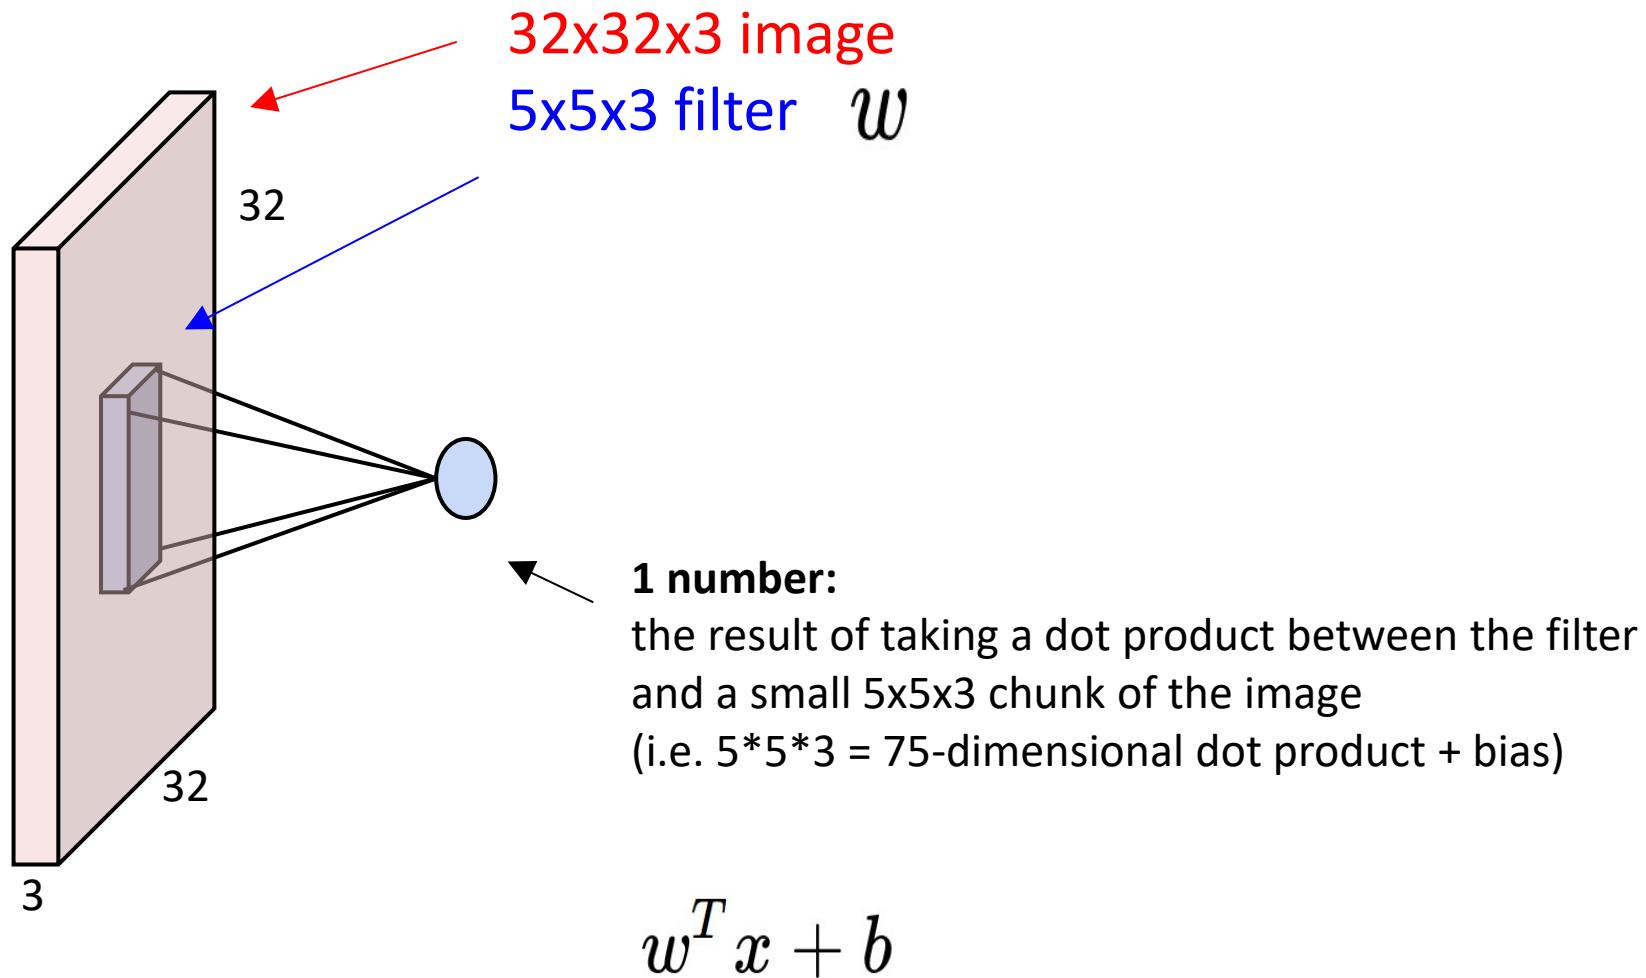


5x5x3 filter

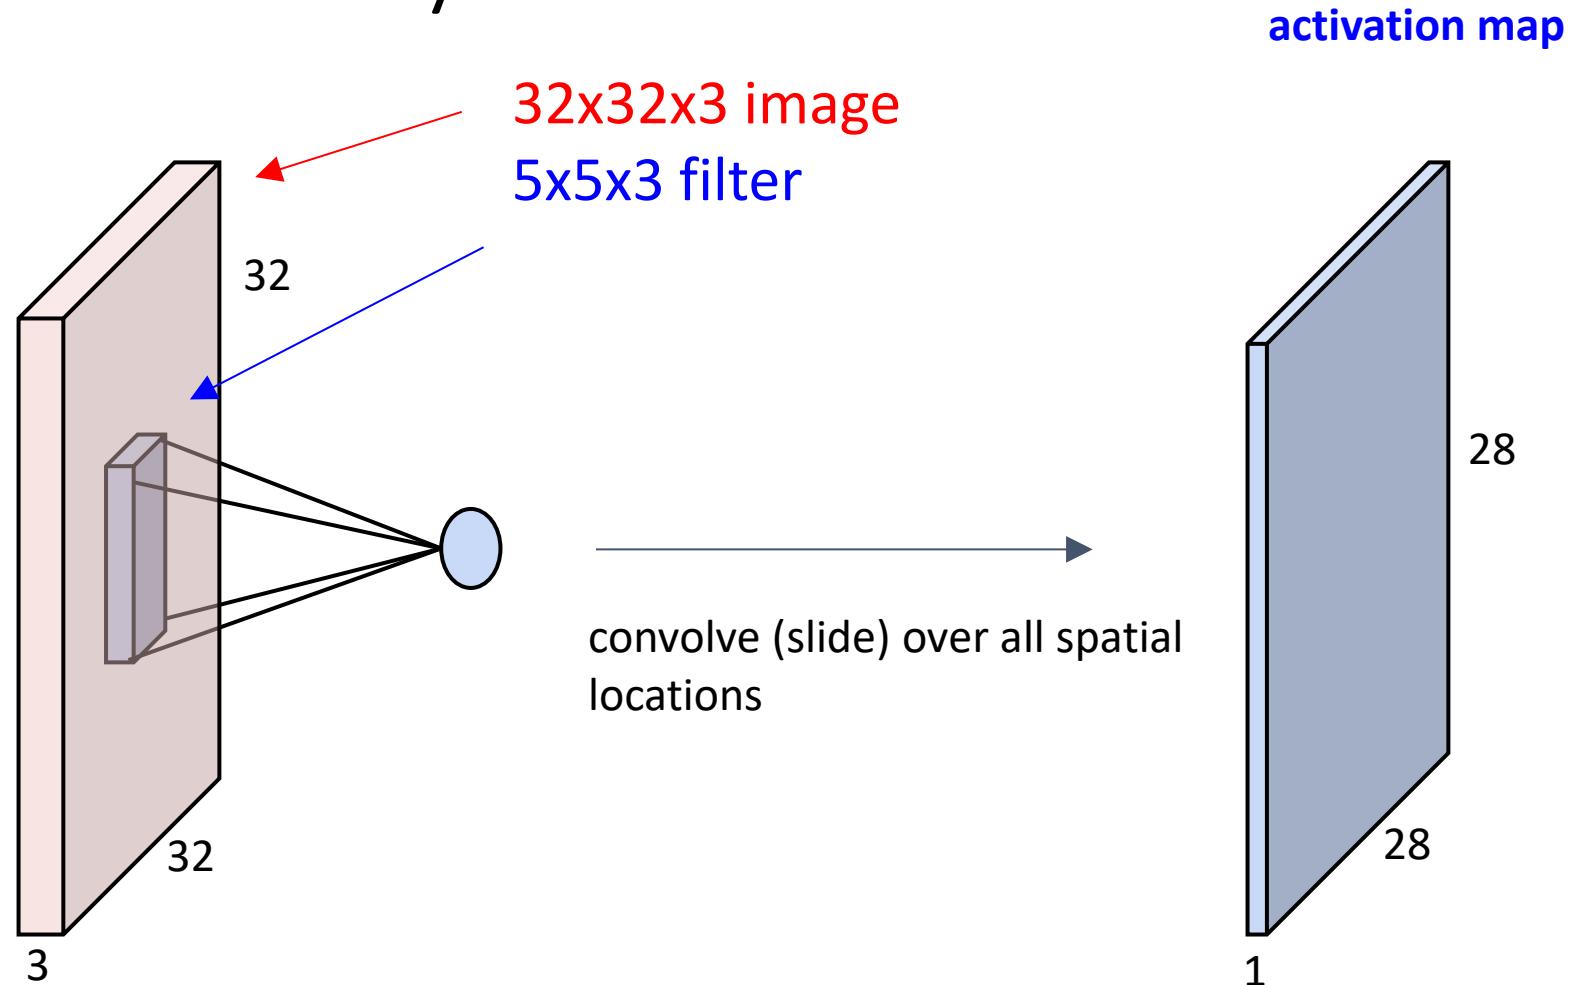


**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

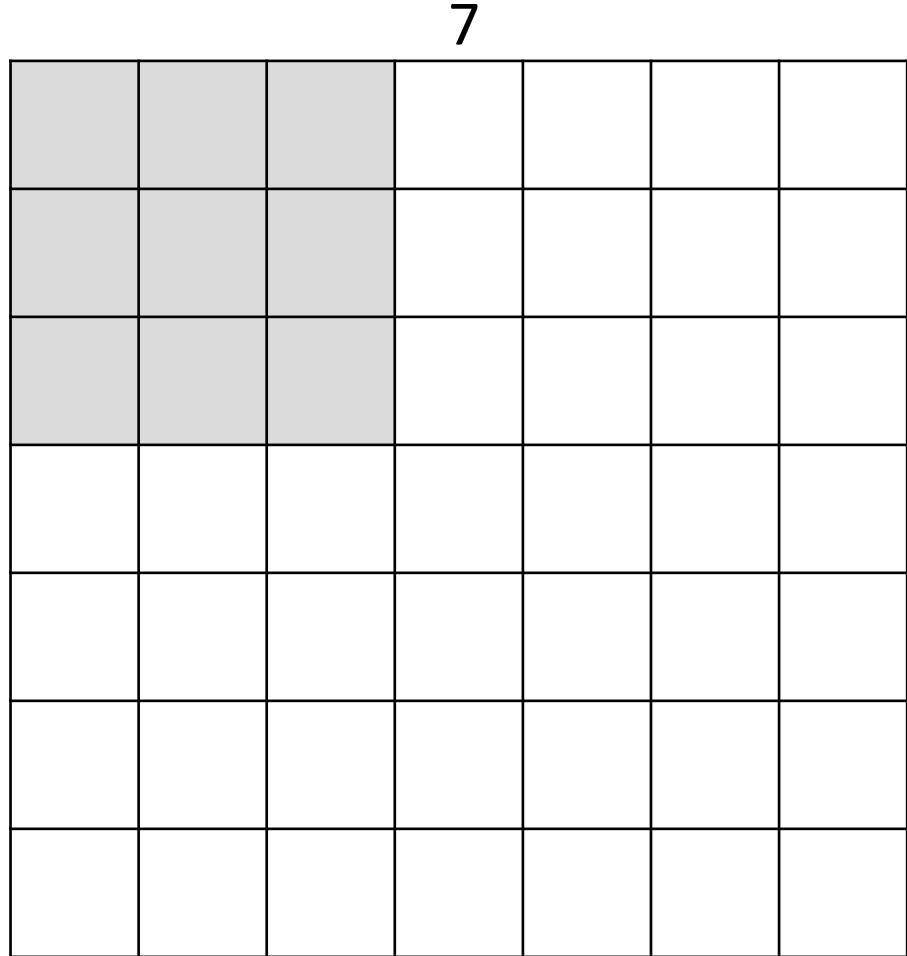
# Convolutional Layer



# Convolutional Layer



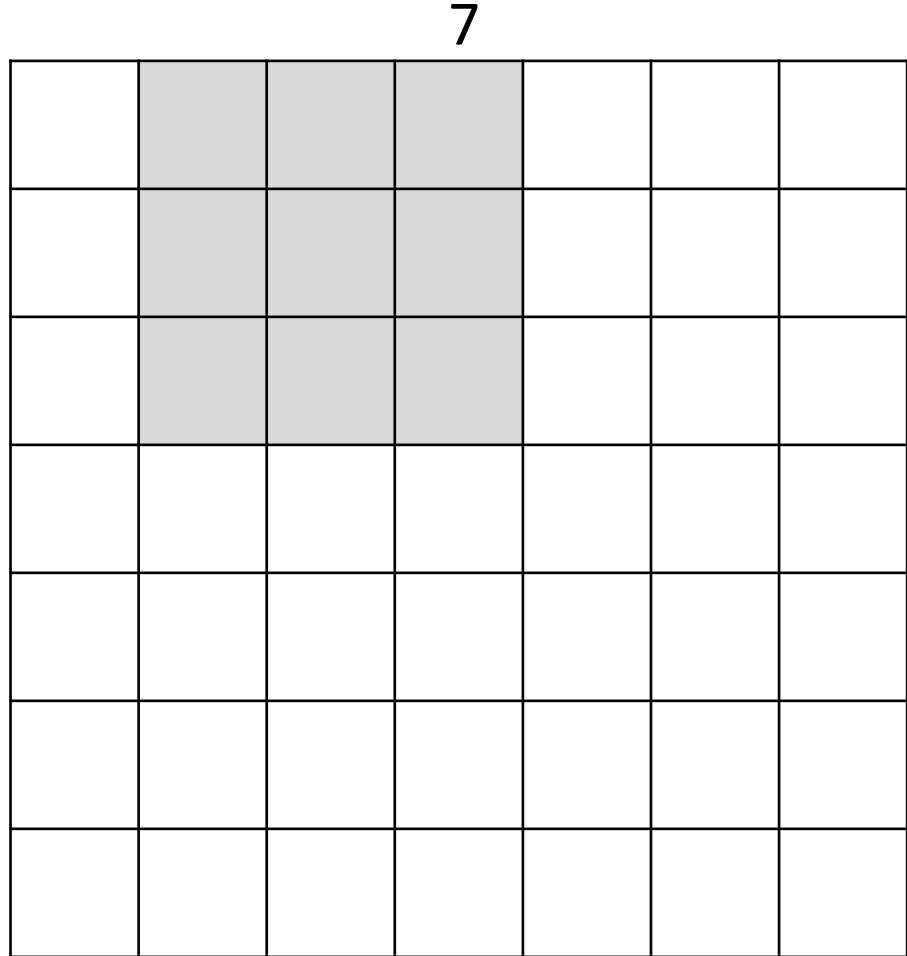
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

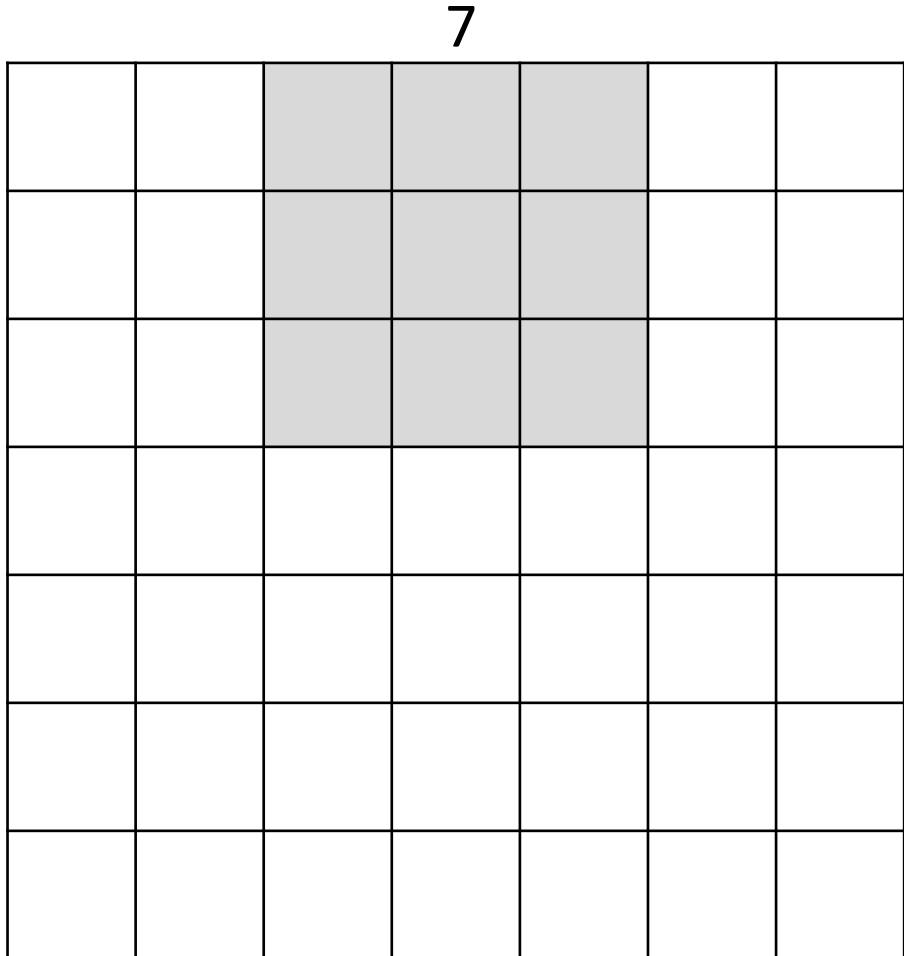
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

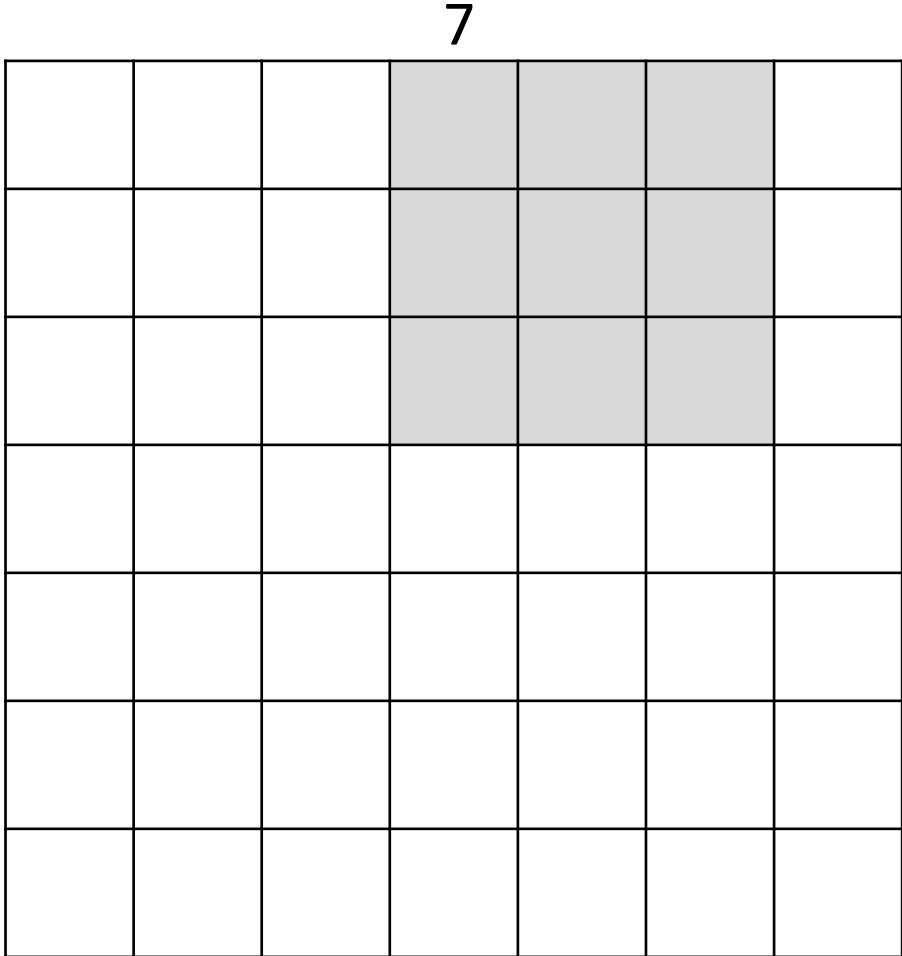
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

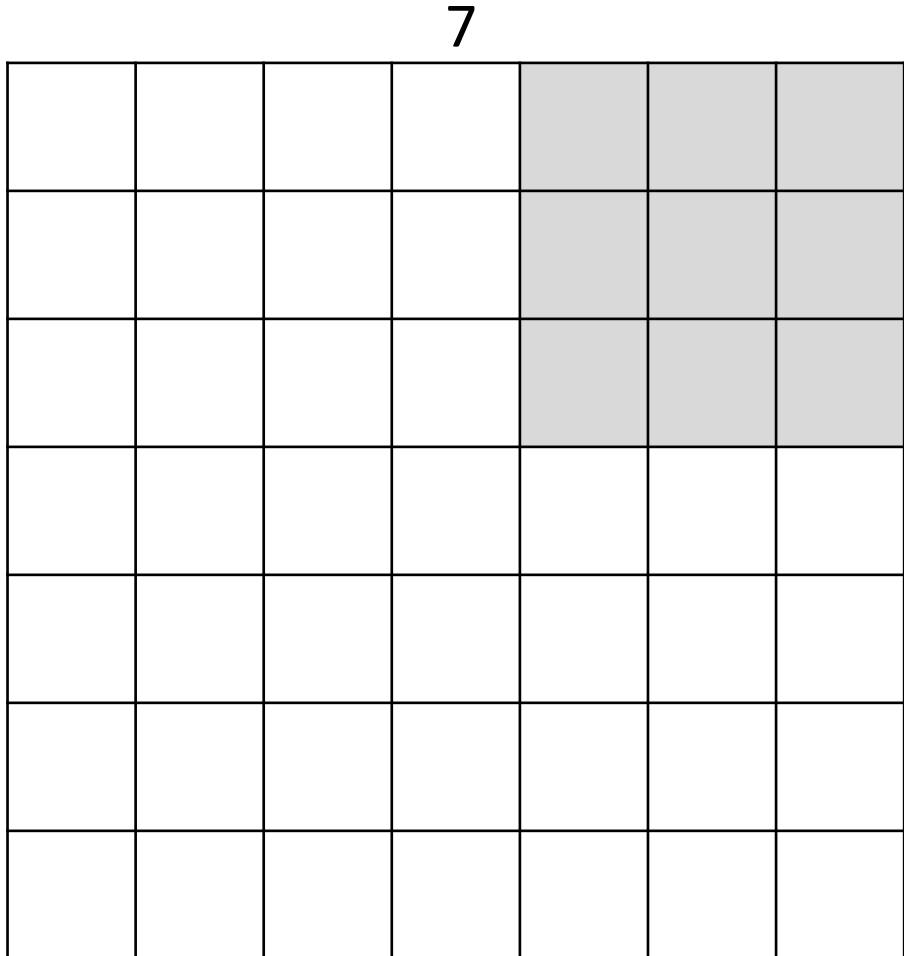
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

# Convolutional Layer

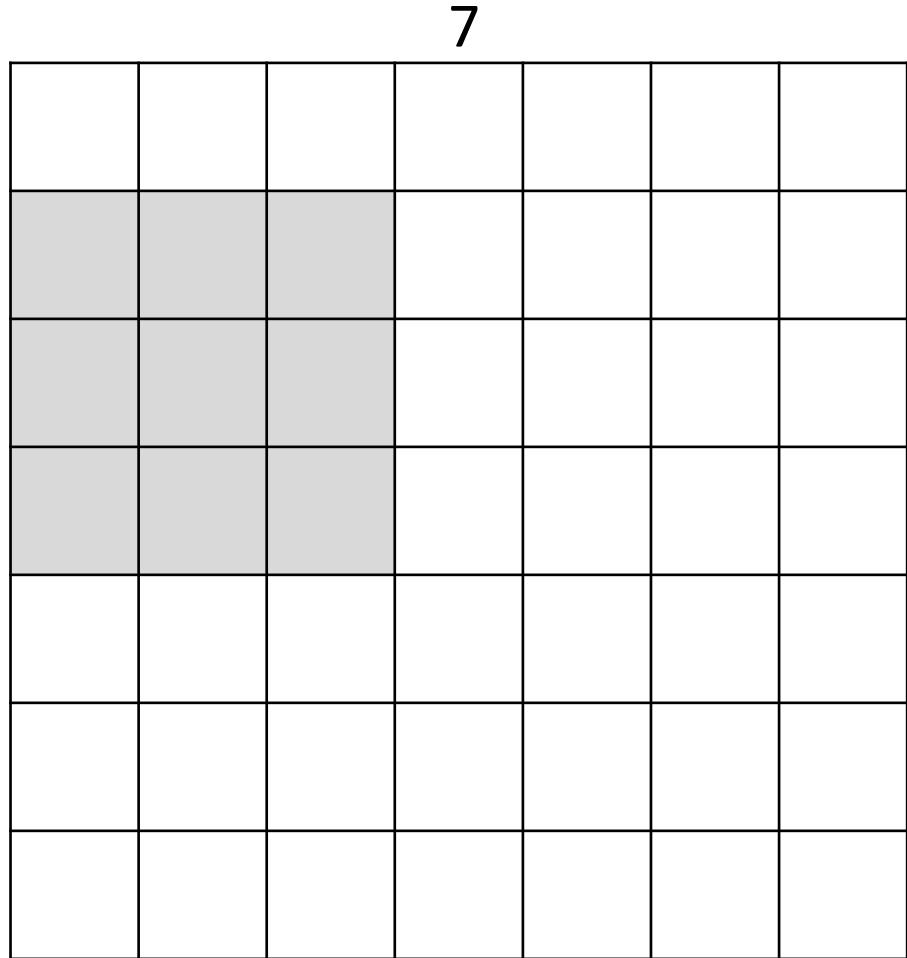


A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

7

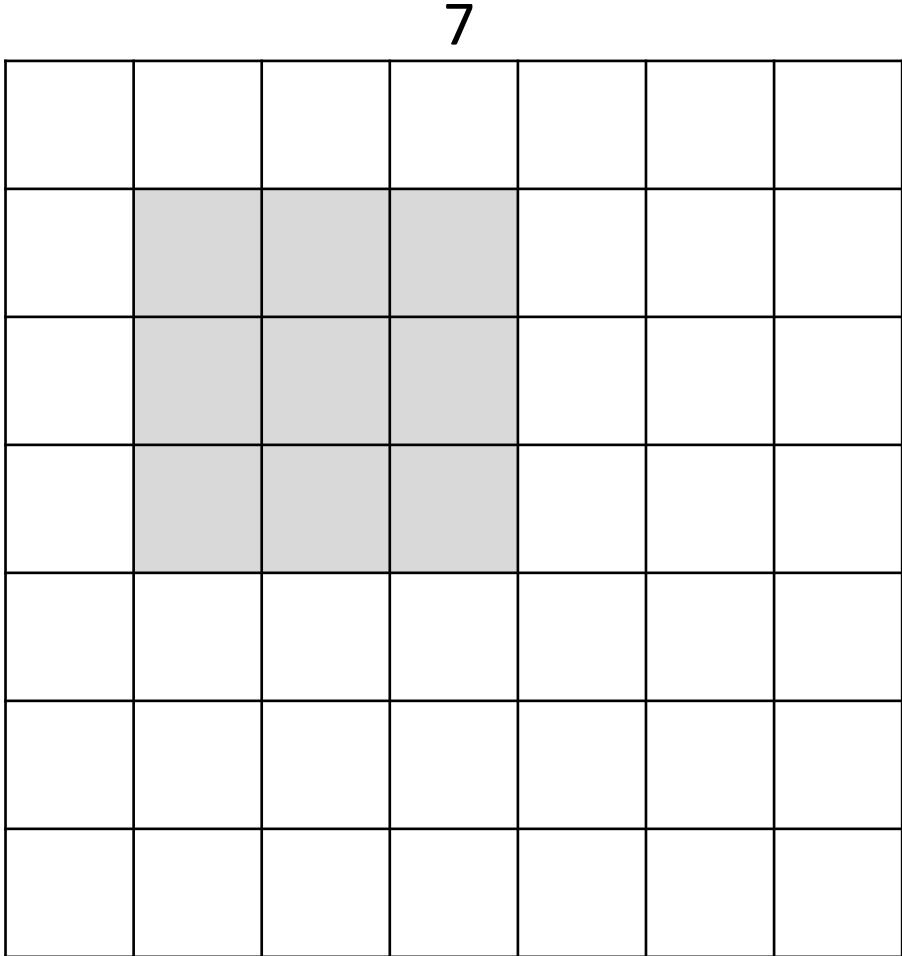
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

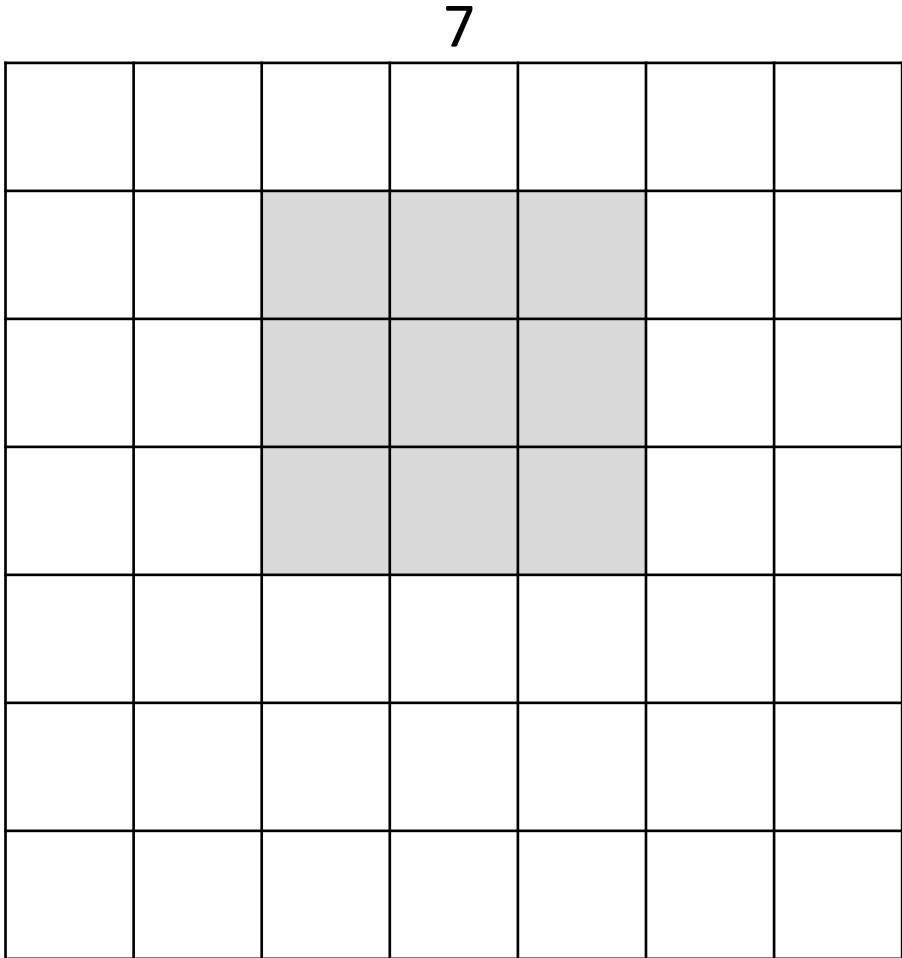
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

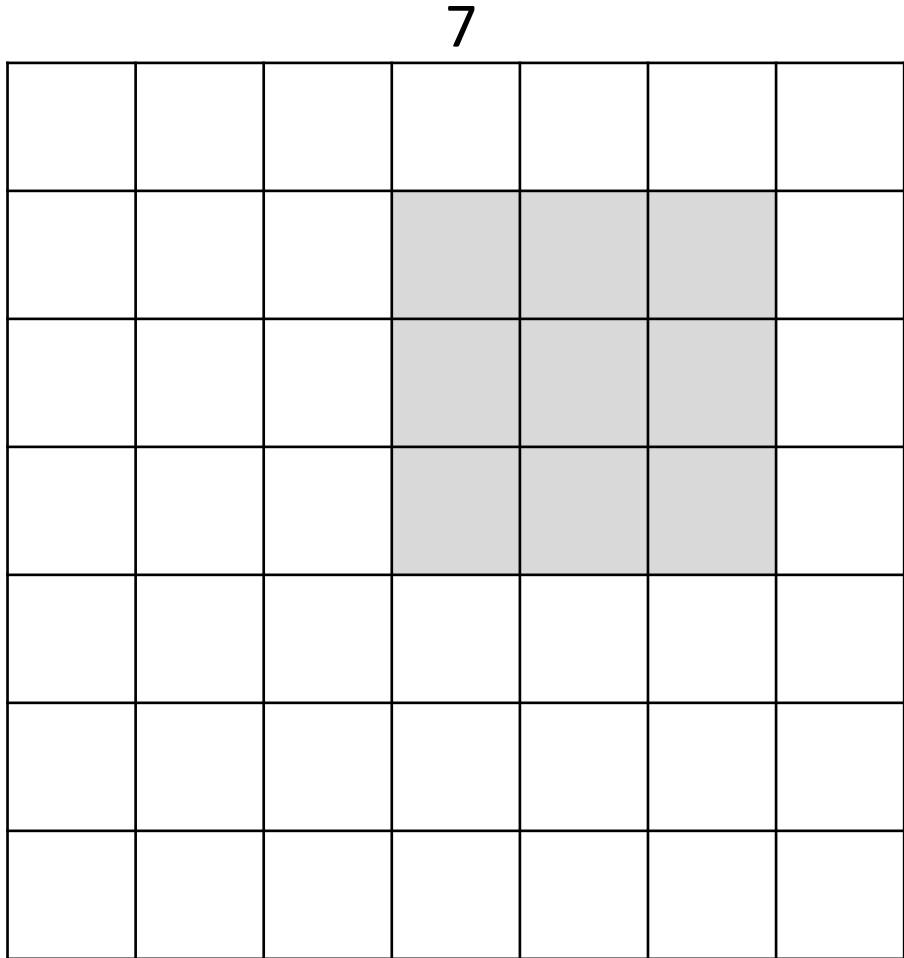
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

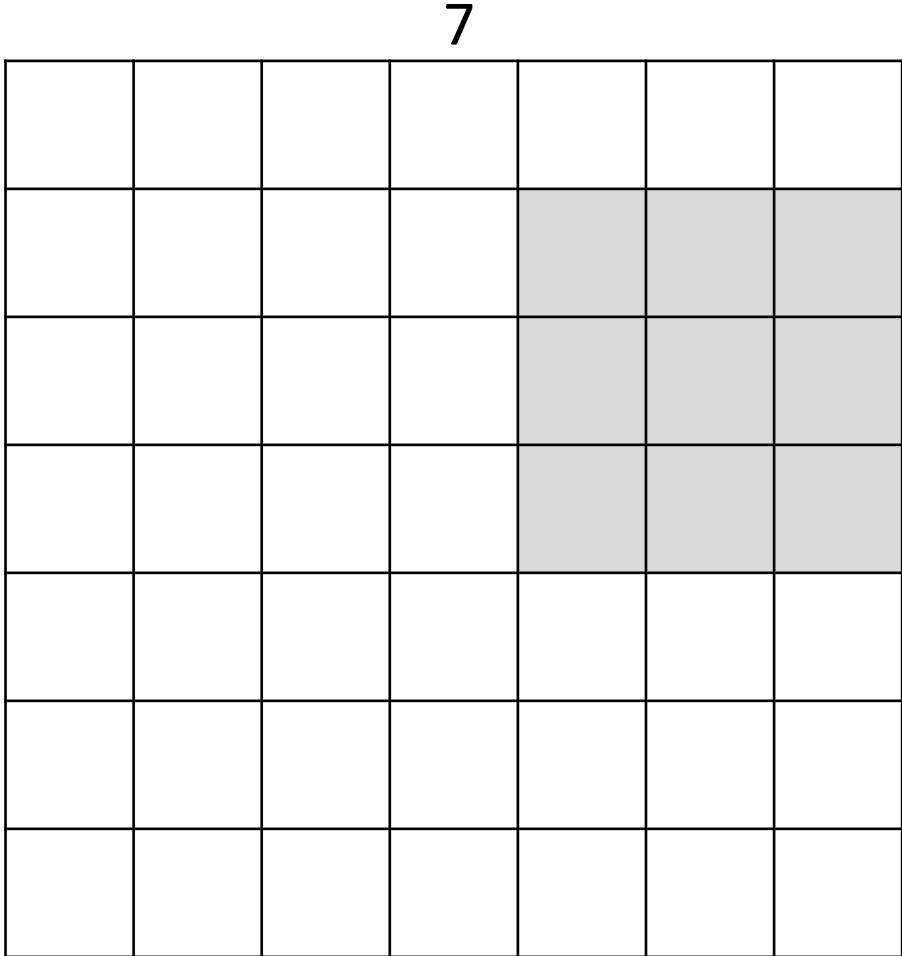
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

# Convolutional Layer

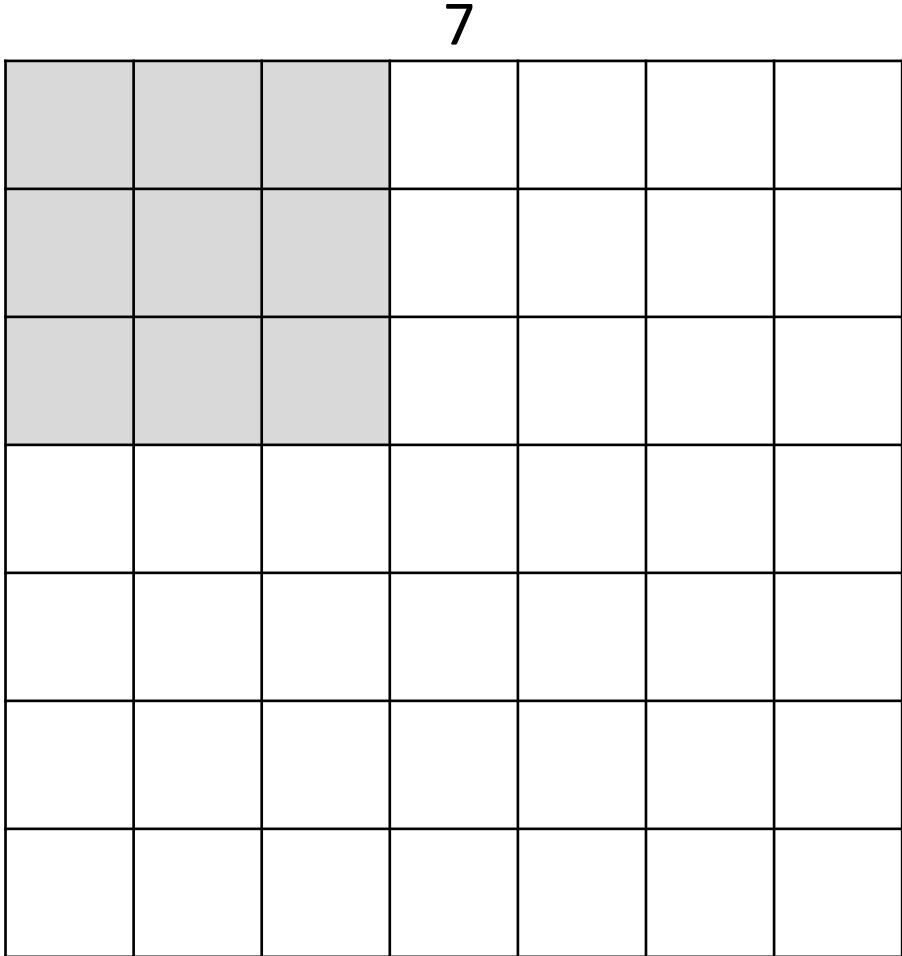


A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, with stride 1

=> **5x5 output**

# Convolutional Layer

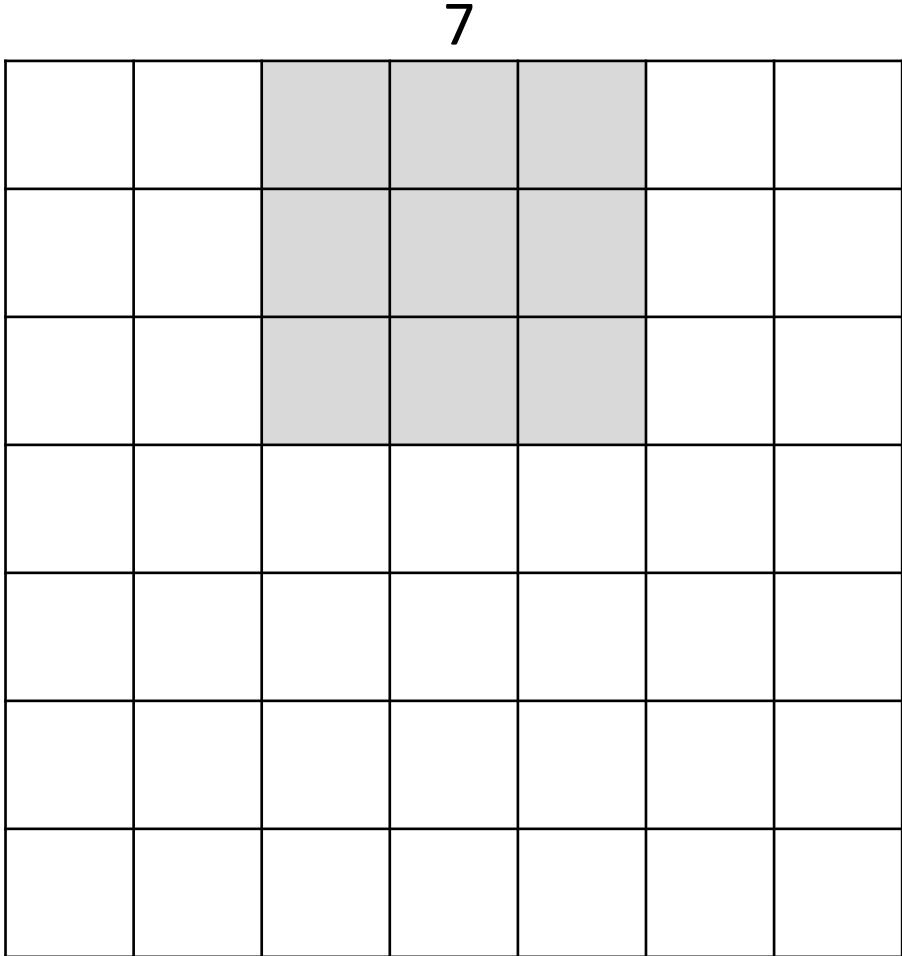


A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, **with stride 2**

7

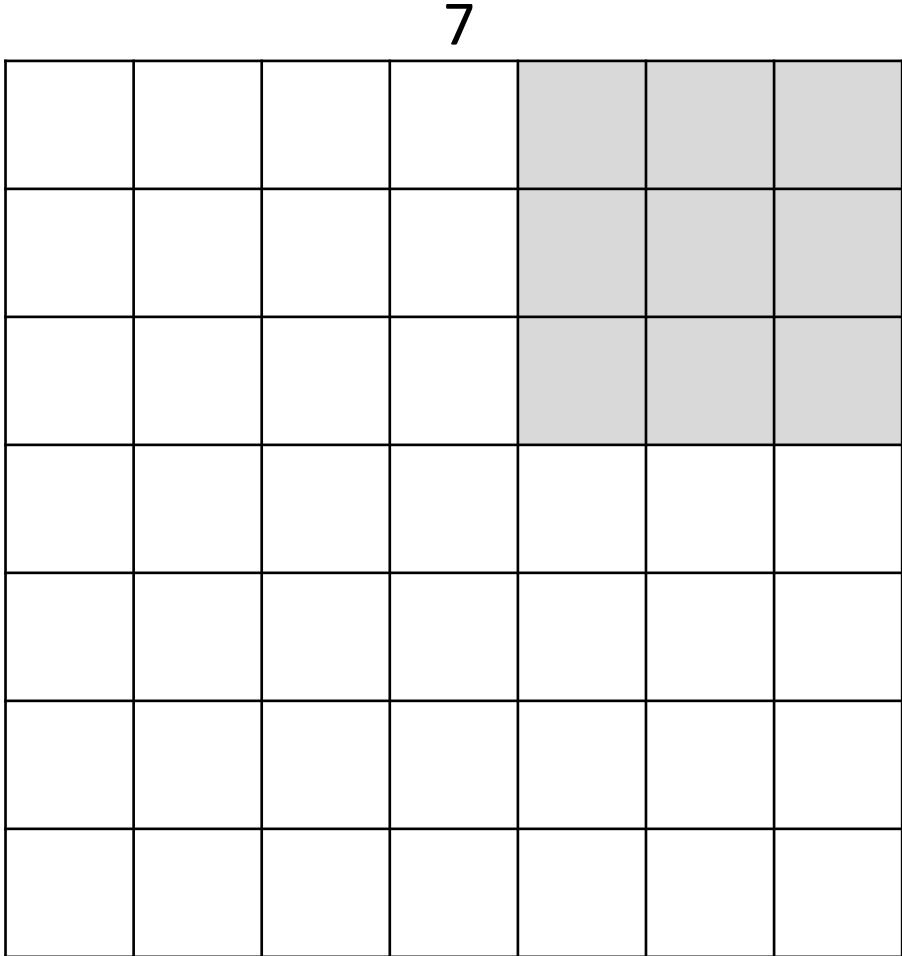
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, **with stride 2**

# Convolutional Layer



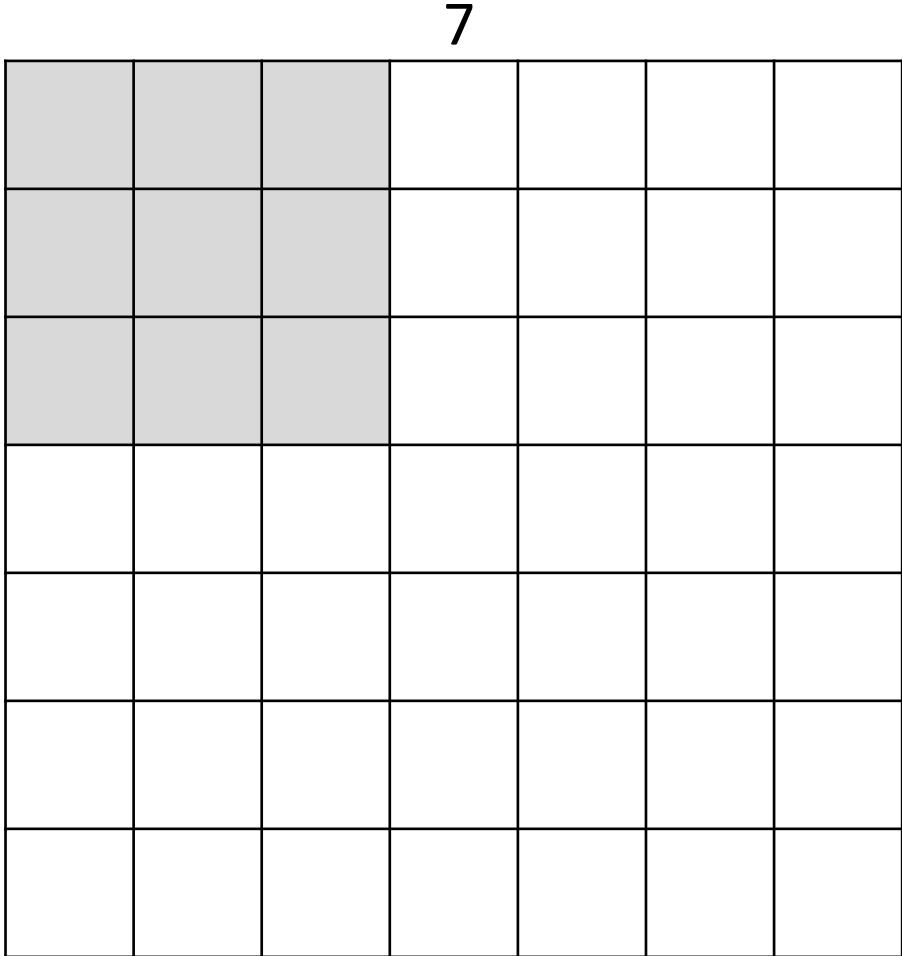
A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, **with stride 2**

=> **3x3 output!**

7

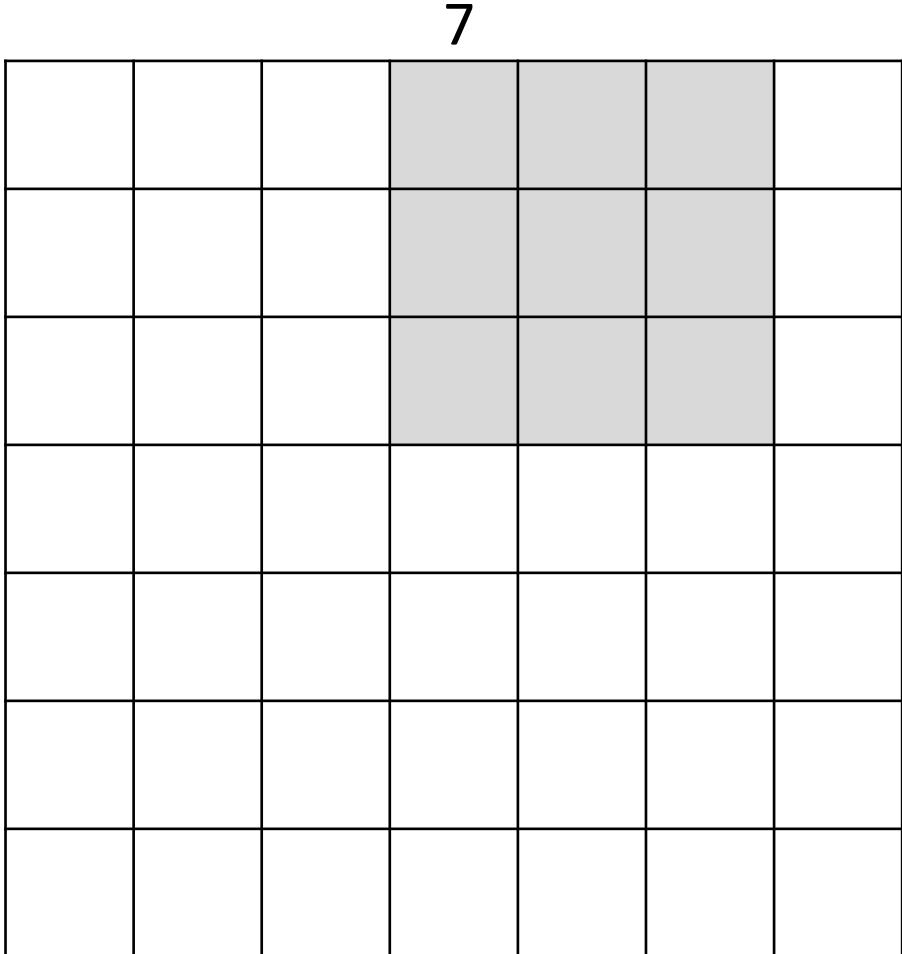
# Convolutional Layer



A closer look at spatial dimensions:

7x7 input (spatially)  
assume 3x3 filter, **with stride 3**

# Convolutional Layer



A closer look at spatial dimensions:

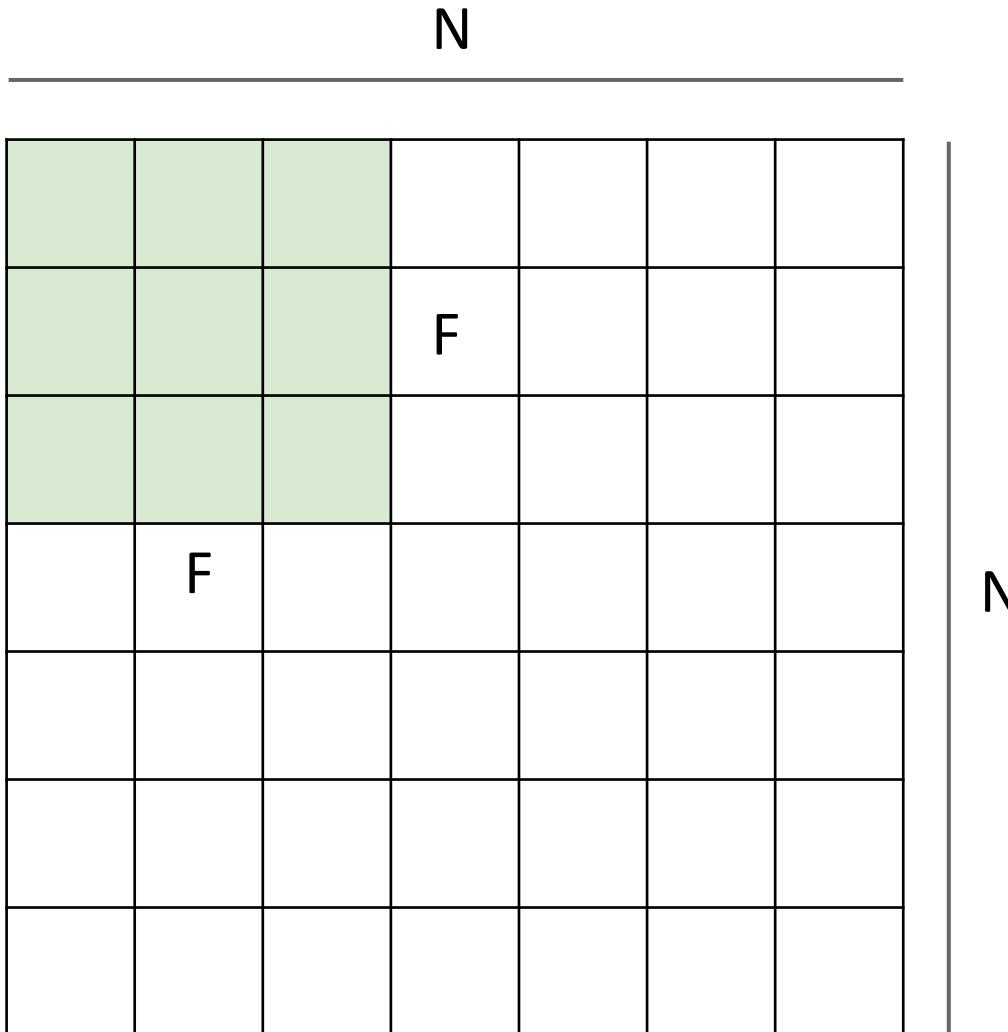
7x7 input (spatially)  
assume 3x3 filter, **with stride 3**

7

**doesn't fit!**

cannot apply 3x3 filter on 7x7  
input with stride 3.

# Convolutional Layer



Output size:  
 **$(N - F) / \text{stride} + 1$**

e.g.  $N = 7, F = 3$ :  
stride 1  $\Rightarrow (7 - 3)/1 + 1 = 5$   
stride 2  $\Rightarrow (7 - 3)/2 + 1 = 3$   
stride 3  $\Rightarrow (7 - 3)/3 + 1 = 2.33$

# Convolutional Layer

0	0	0	0	0	0			
0								
0								
0								
0								

In practice: Common to zero pad the border

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

(recall:)

**(N - F) / stride + 1**

# Convolutional Layer

0	0	0	0	0	0			
0								
0								
0								
0								

In practice: Common to zero pad the border

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

**7x7 output!**

# Convolutional Layer

0	0	0	0	0	0			
0								
0								
0								
0								

In practice: Common to zero pad the border

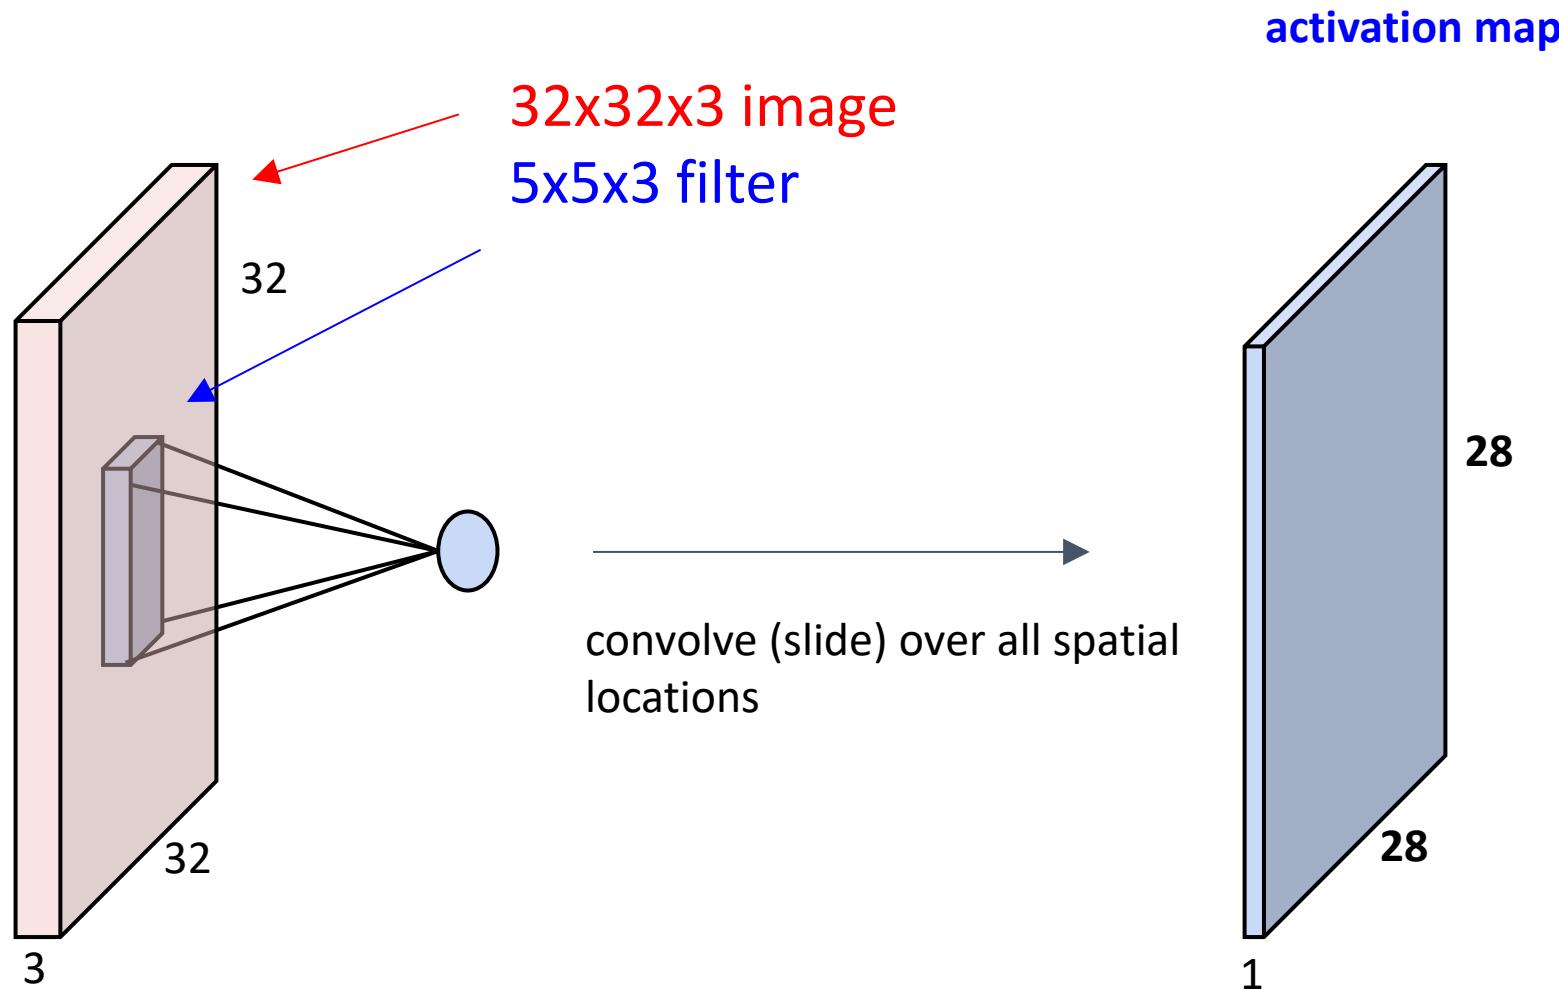
in general, common to see CONV layers with stride 1, filters of size  $F \times F$ , and zero-padding with  $(F-1)/2$ . (will preserve size spatially)

e.g.  $F = 3 \Rightarrow$  zero pad with 1

$F = 5 \Rightarrow$  zero pad with 2

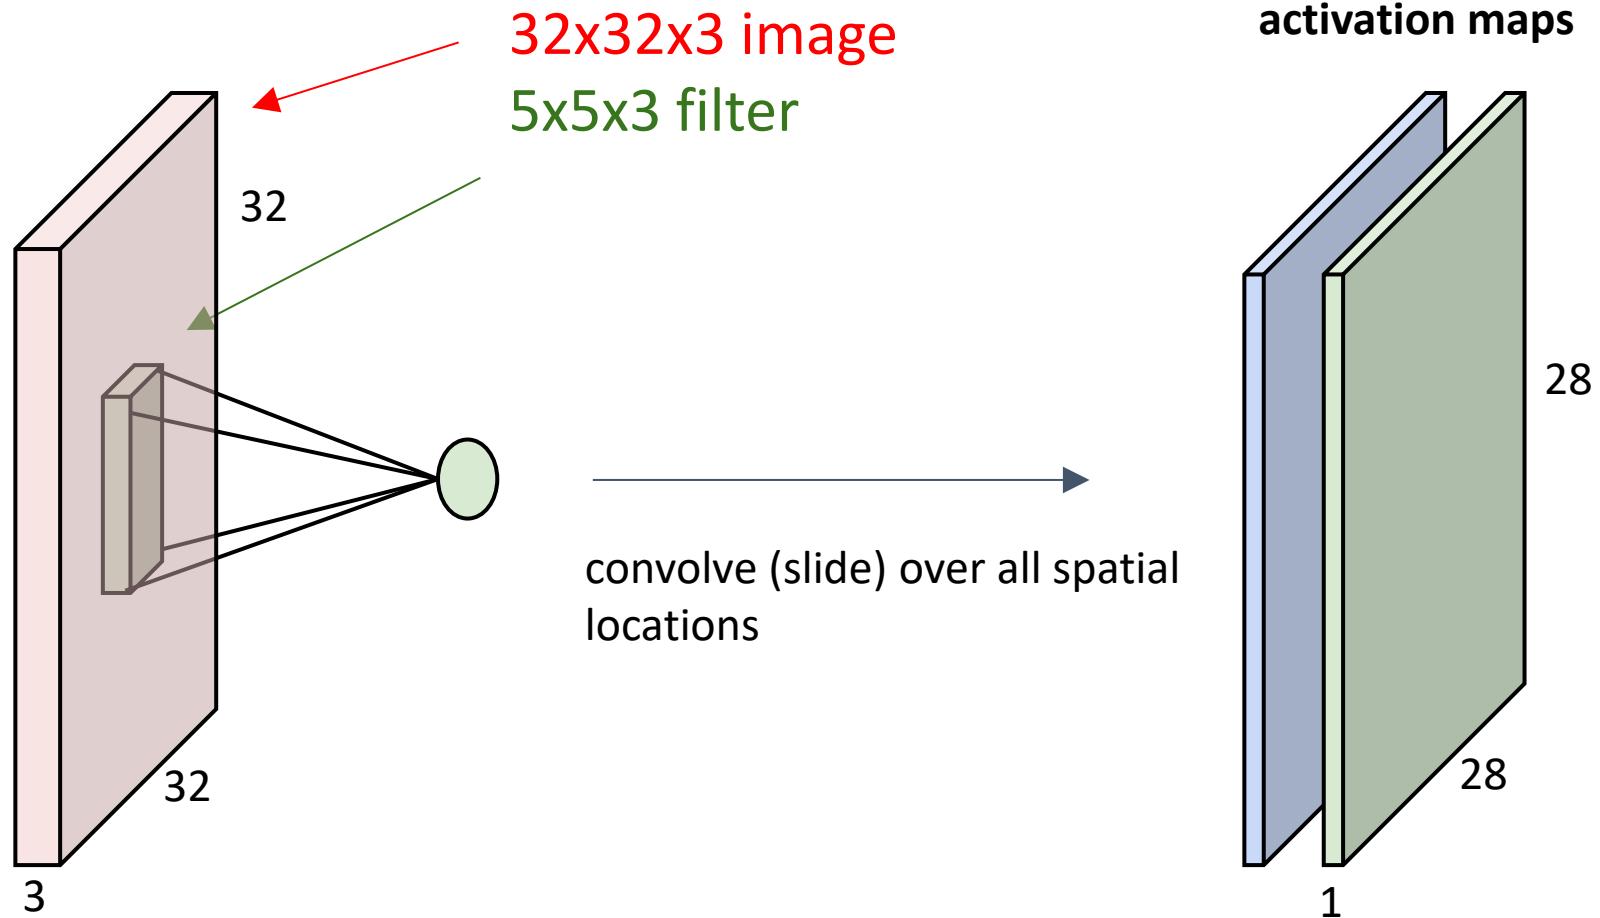
$F = 7 \Rightarrow$  zero pad with 3

# Convolutional Layer

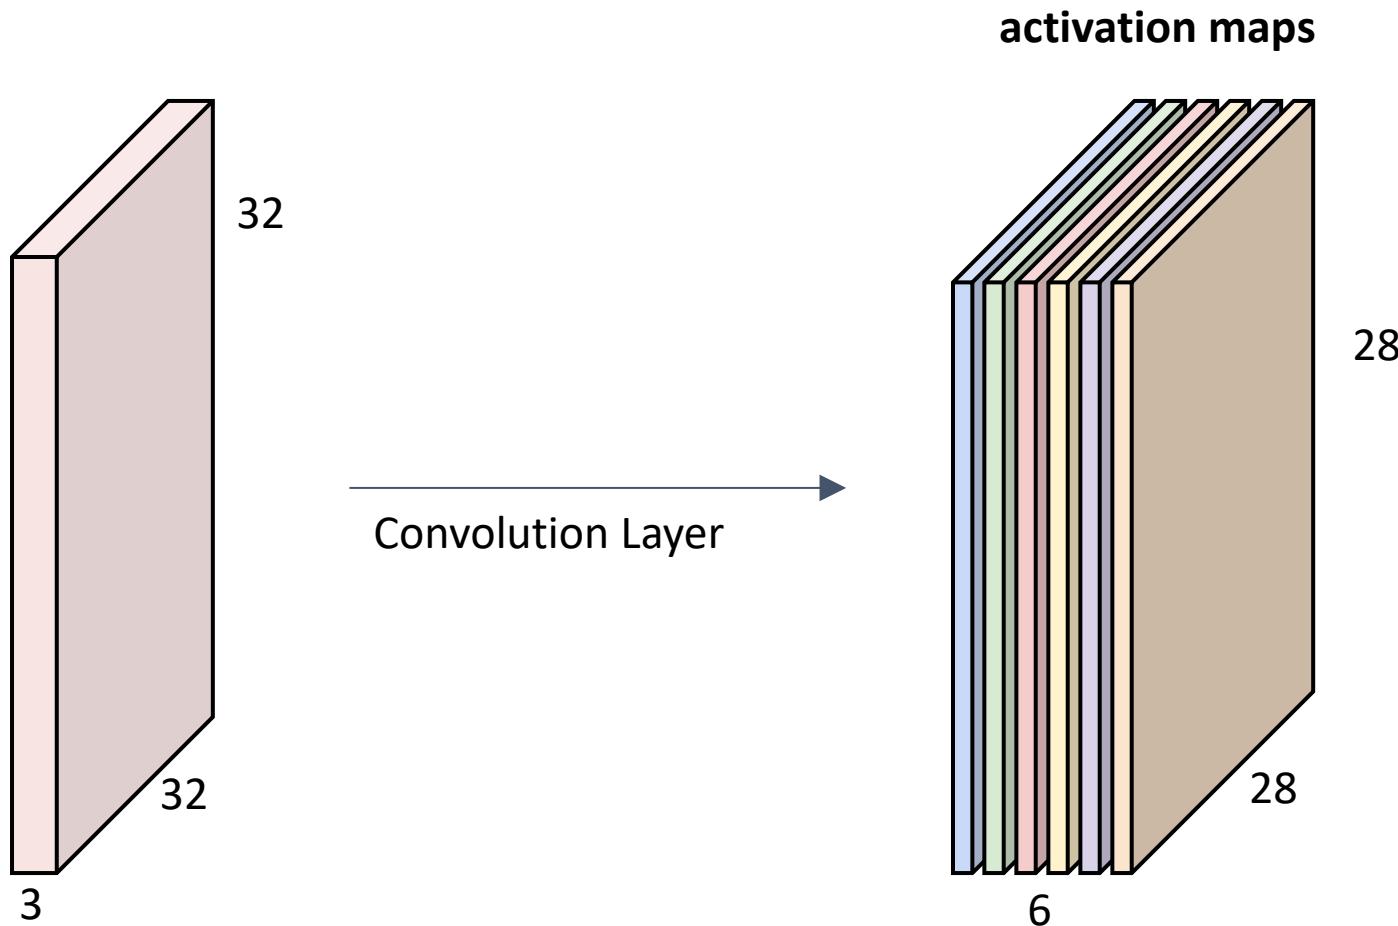


# Convolutional Layer

consider a second, **green** filter



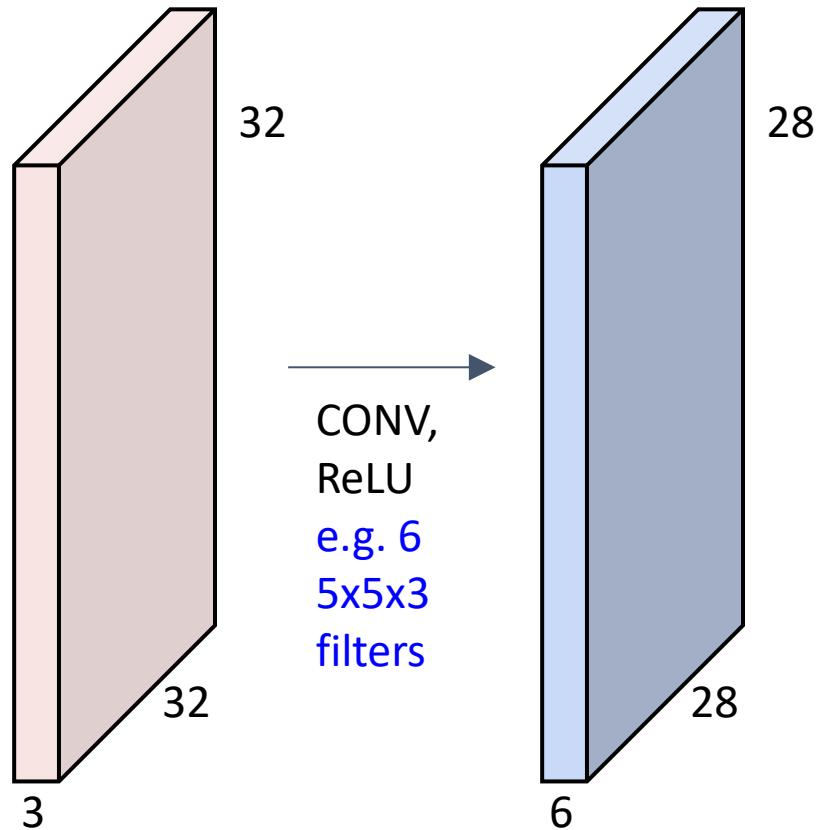
# Convolutional Layer



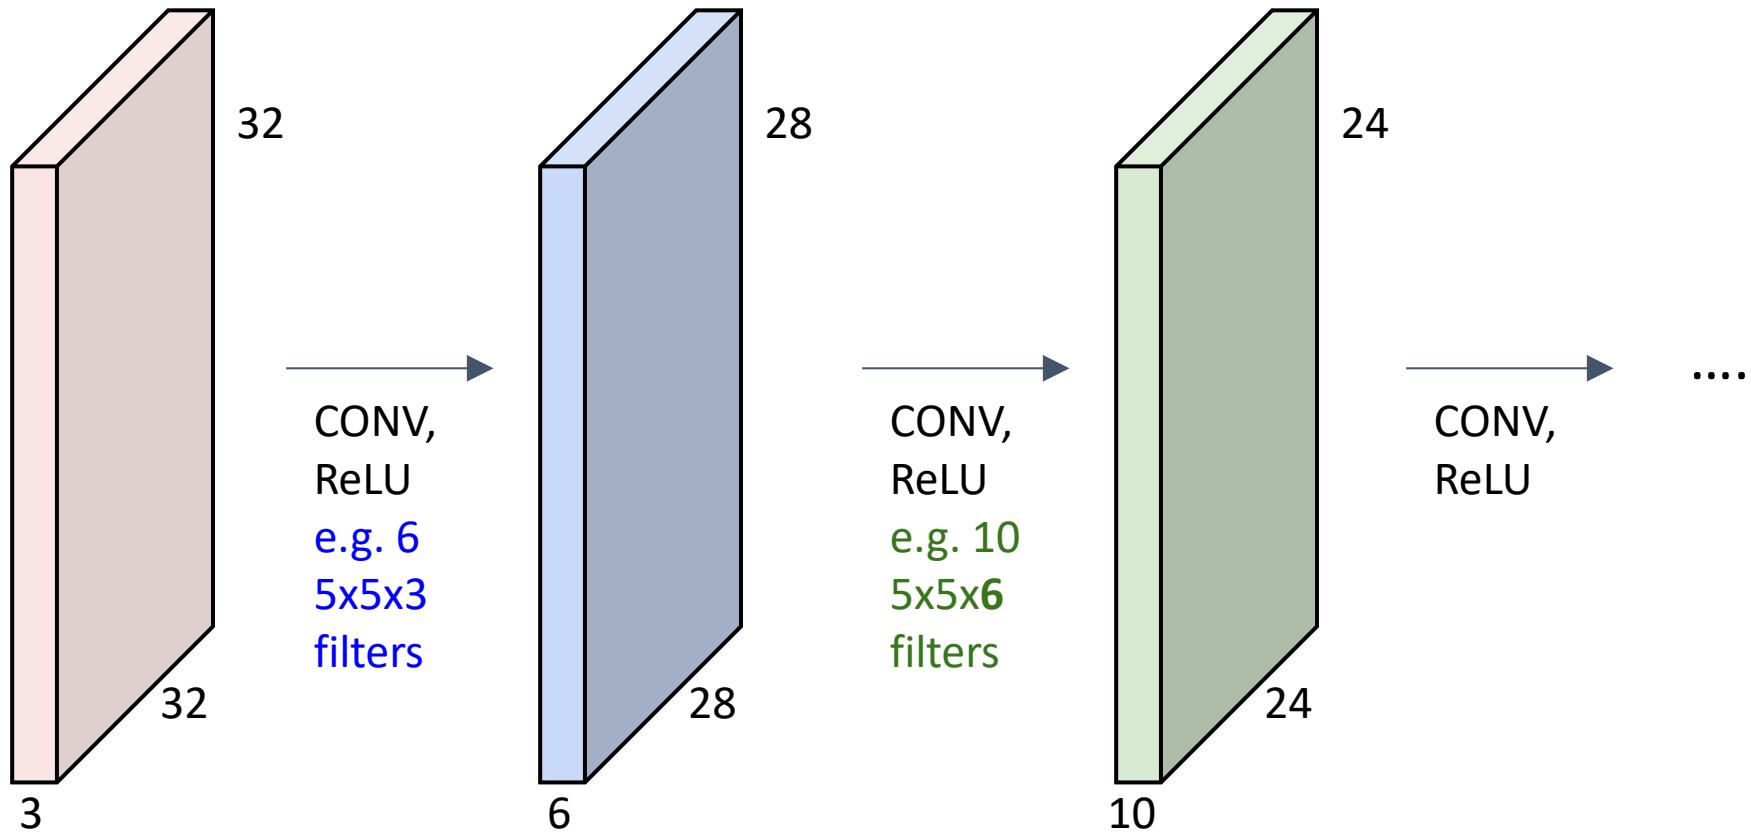
For example, if we had 6 5x5 filters, we'll get 6 separate activation maps

We stack these up to get a “new image” of size 28x28x6!

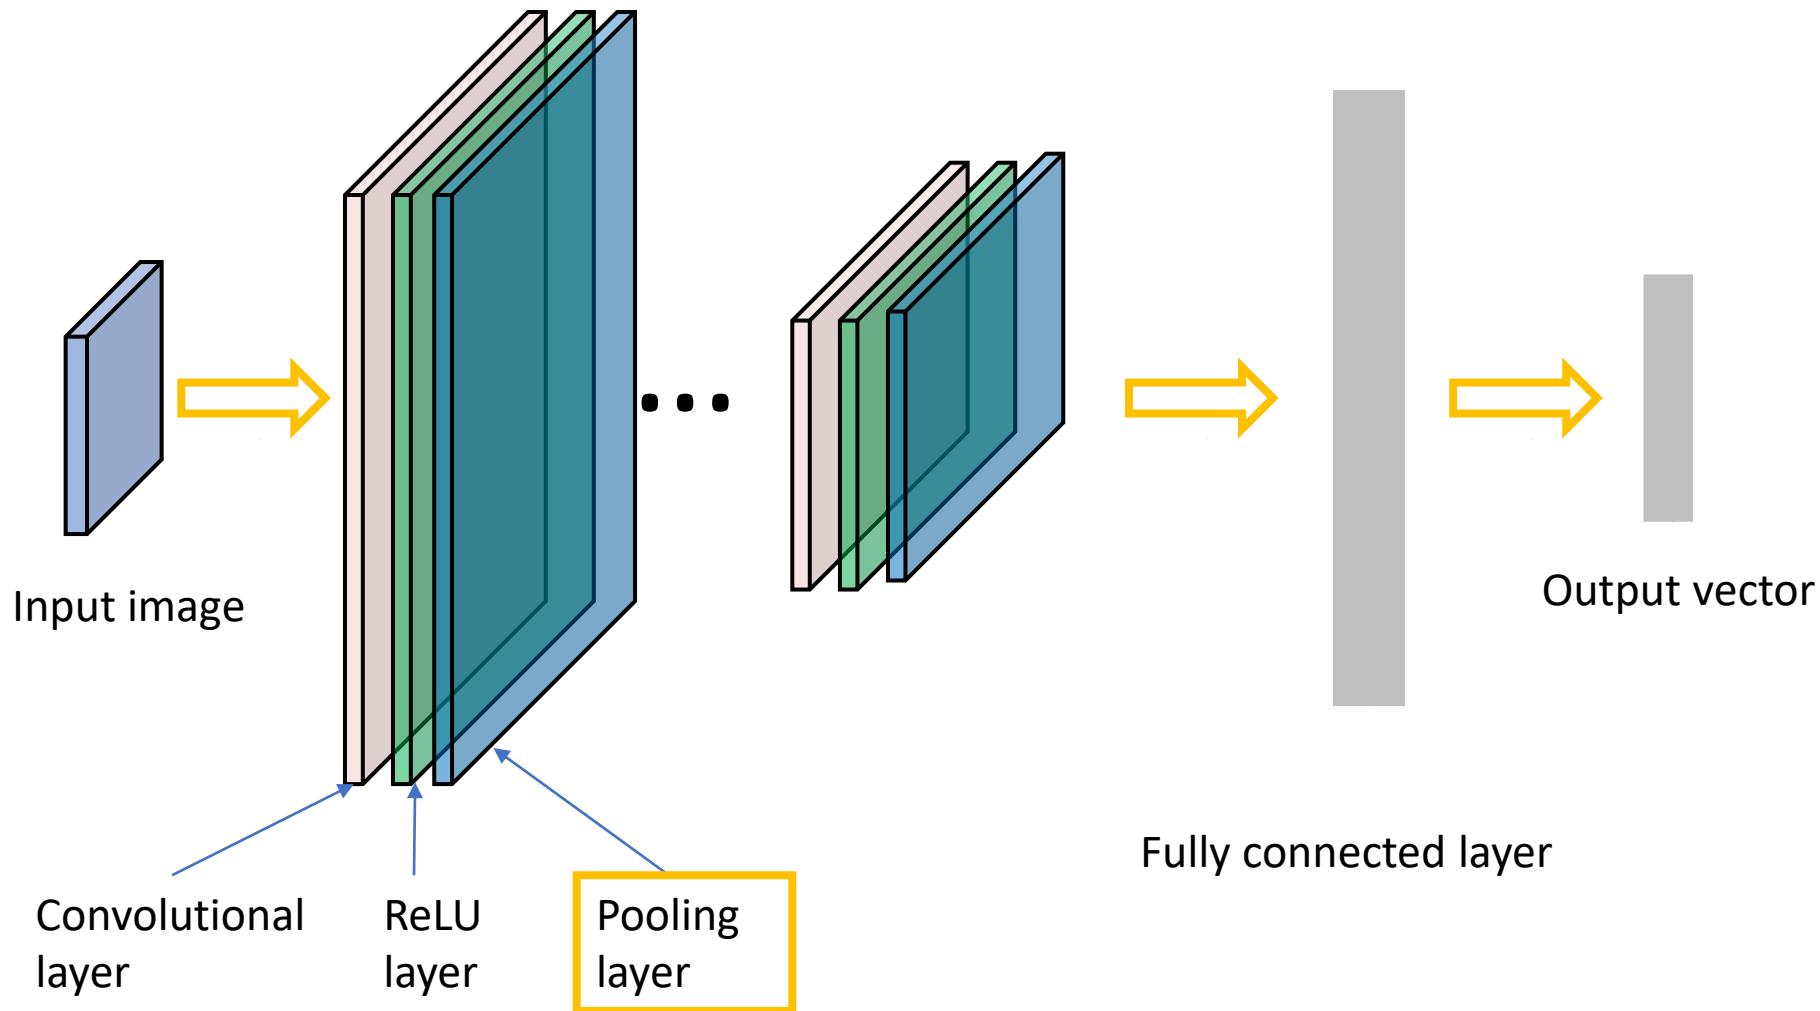
**Preview:** ConvNet is a sequence of Convolution Layers, interspersed with activation functions



**Preview:** ConvNet is a sequence of Convolutional Layers, interspersed with activation functions

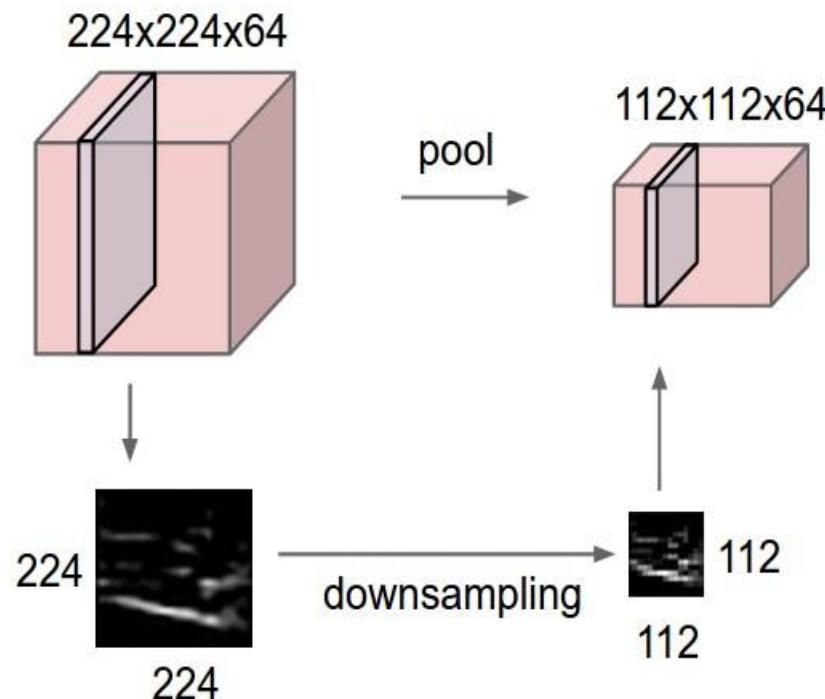


# Convolutional Neural Networks

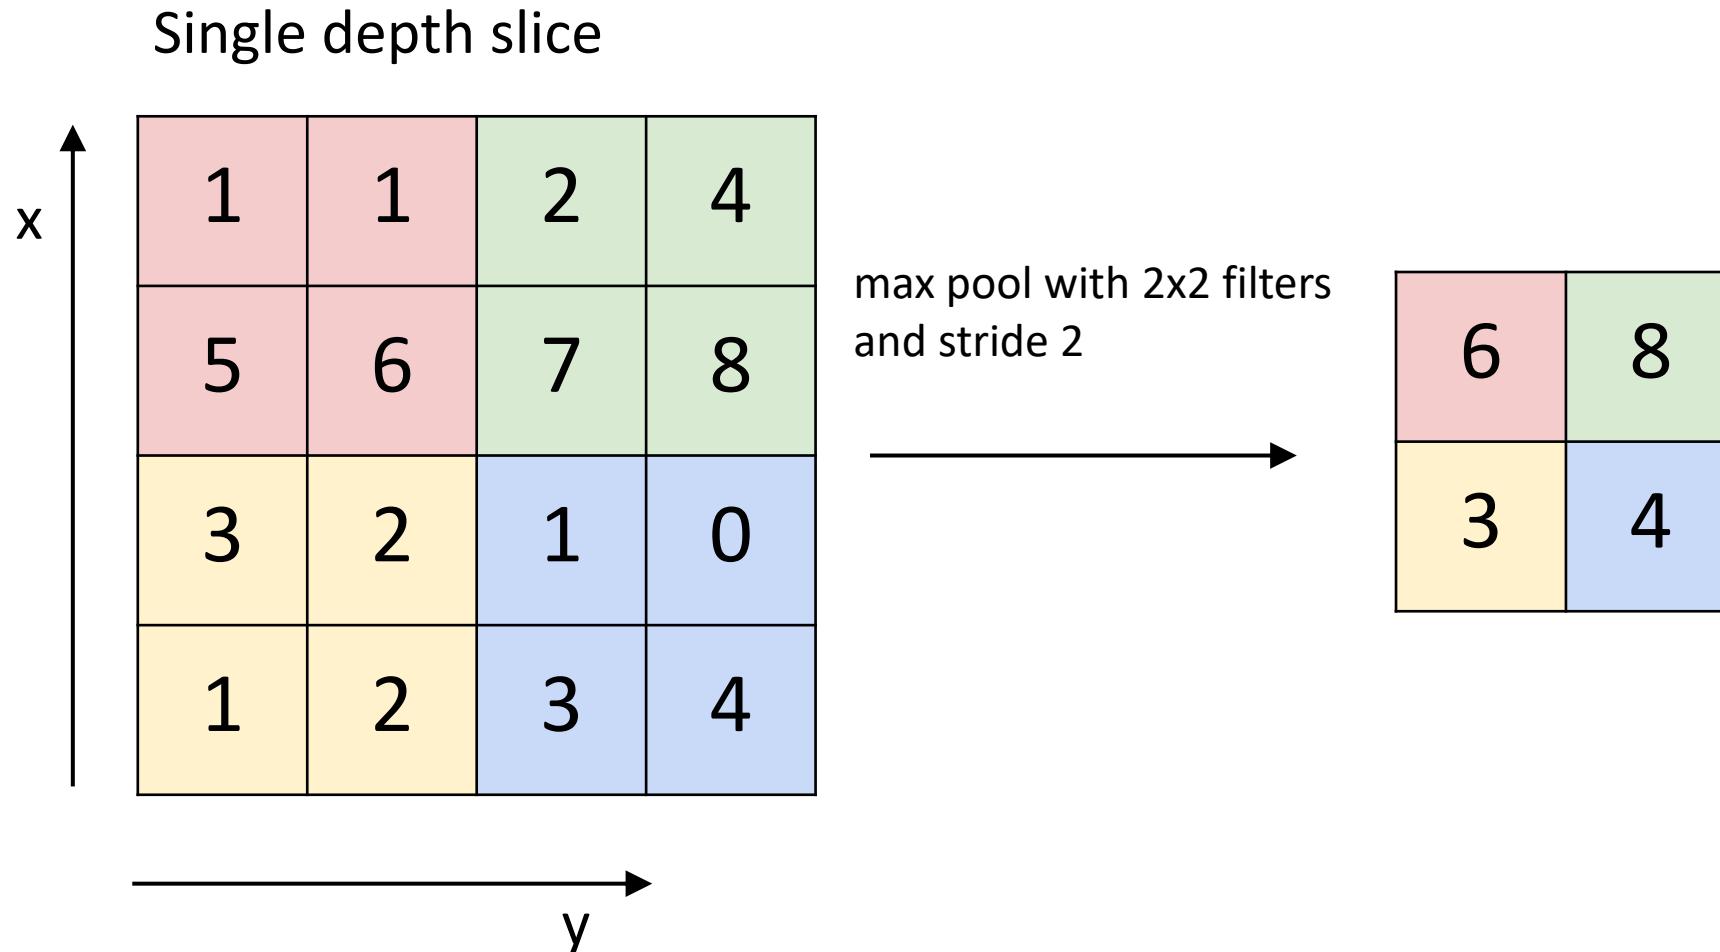


# Pooling Layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



# Max Pooling Layer



# My Python Setup

# Further Reading

- Stanford CS231n, lecture 5, Convolutional Neural Networks  
<http://cs231n.stanford.edu/schedule.html>
- Deep learning with PyTorch  
[https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)
- AlexNet (2012):  
<https://papers.nips.cc/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html>
- Vgg16 (2014): <https://arxiv.org/abs/1409.1556>
- GoogleNet (2014): <https://arxiv.org/abs/1409.4842>
- ResNet (2015): <https://arxiv.org/abs/1512.03385>