



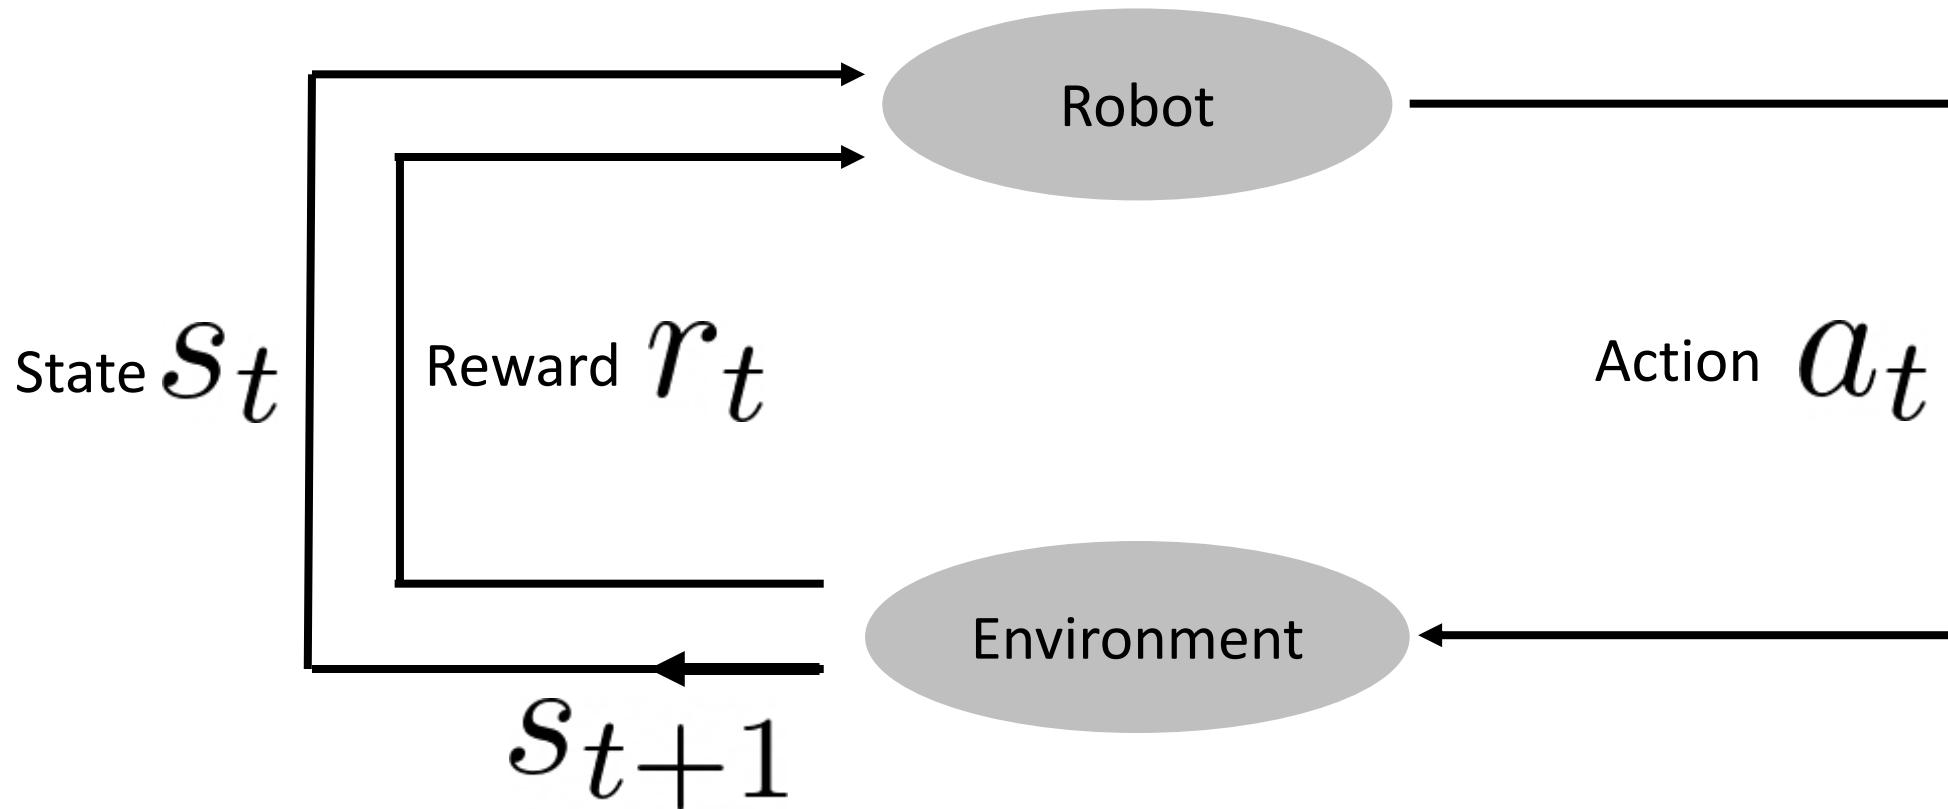
# Reinforcement Learning: Overview and Foundations

CS 6341 Robotics

Professor Yu Xiang

The University of Texas at Dallas

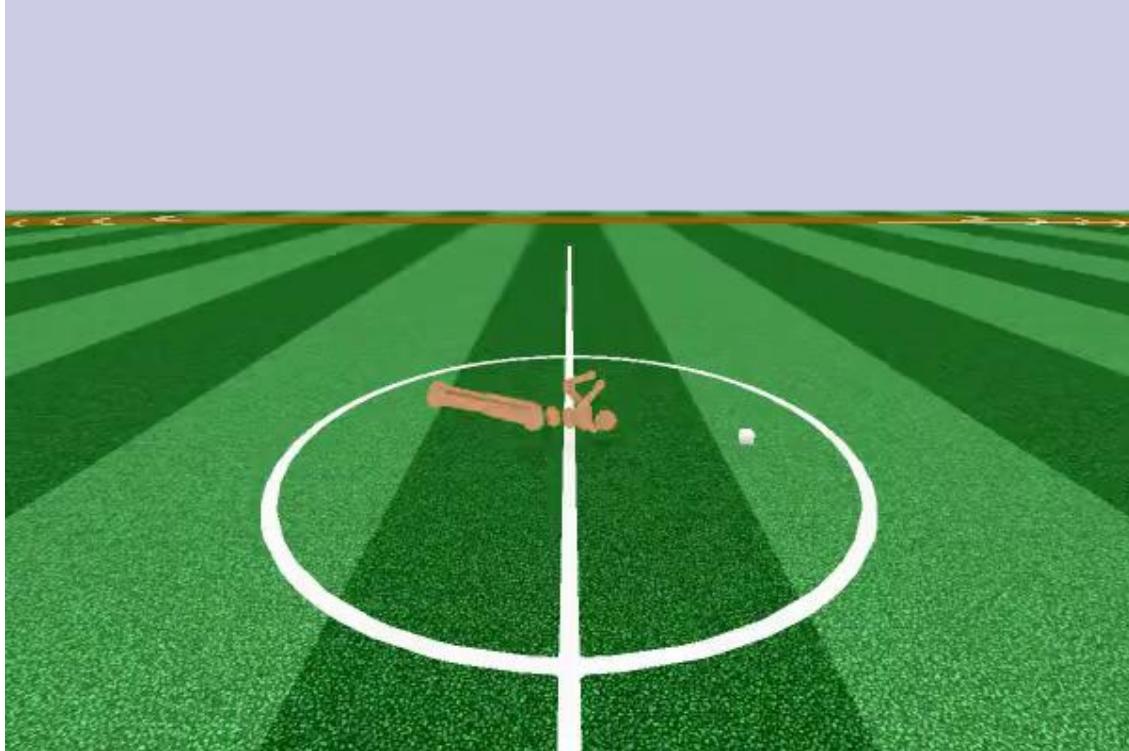
# Reinforcement Learning



Reinforcement Learning:  
Imitation Learning:

$$a_t = \pi(s_t)$$

# RL Examples



Control

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro.html)

# RL Examples

## Unitree RL GYM

English | [cn中文](#)

This is a repository for reinforcement learning implementation based on Unitree robots, supporting Unitree Go2, H1, H1\_2, and G1.

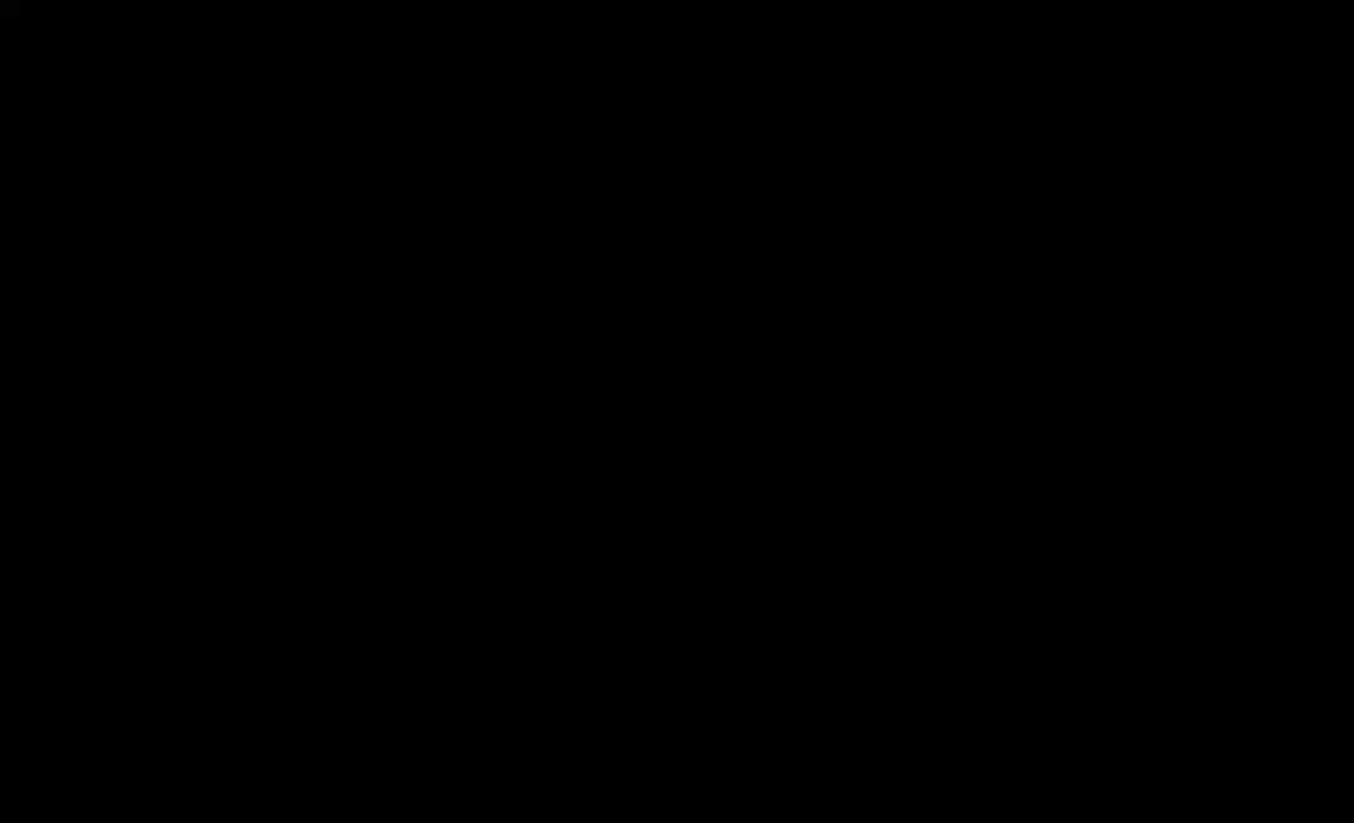
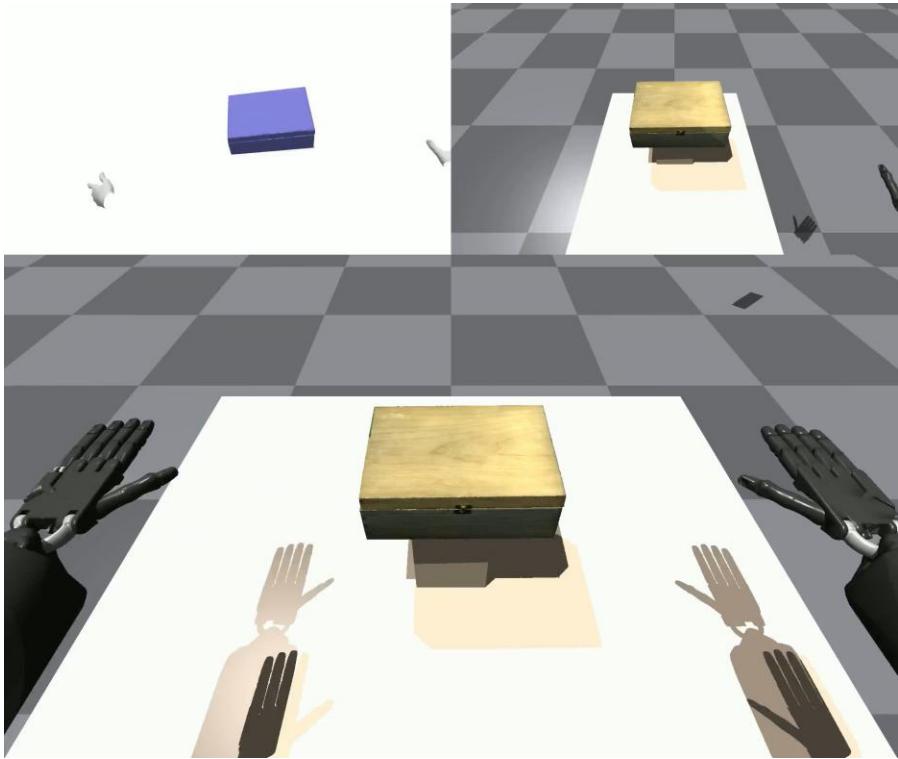


[https://github.com/unitreerobotics/unitree\\_rl\\_gym](https://github.com/unitreerobotics/unitree_rl_gym)



Imitation learning & reinforcement learning

# RL Examples



<https://cypypccpy.github.io/obj-dex.github.io/>

<https://playground.mujoco.org/>

# RL Concepts

- State  $S$  : a complete description of the state of the world
- Observation  $O$ : partial description of a state
  - Fully observed vs. partially observed
  - For example: images
- Action space: the set of all valid actions in a given environment
  - Discrete action space vs. continuous action space  $a$
- Policies: a policy is a rule used by an agent to decide what action to take
  - Deterministic policy  $a_t = \mu(s_t)$
  - Stochastic policy  $a_t \sim \pi(\cdot | s_t)$

# RL Concepts

- Parameterized policies

$$a_t = \mu_\theta(s_t)$$

$$a_t \sim \pi_\theta(\cdot | s_t)$$

k dimensional action

- Joint position
- Gripper pose
- Joint velocity, etc.

- Deterministic policy

```
pi_net = nn.Sequential(  
    nn.Linear(obs_dim, 64),  
    nn.Tanh(),  
    nn.Linear(64, 64),  
    nn.Tanh(),  
    nn.Linear(64, act_dim)  
)
```

- Stochastic policy

- Categorical policy for discrete actions

$$\log \pi_\theta(a|s) = \log [P_\theta(s)]_a$$

- Diagonal Gaussian policy: mean action

$$\mu_\theta(s)$$

Log standard deviation

$$\log \sigma_\theta(s)$$

$(-\infty, \infty)$

# RL Concepts

- Diagonal Gaussian policy

- Sampling  $a = \mu_\theta(s) + \sigma_\theta(s) \odot z$   $z \sim \mathcal{N}(0, I)$  k dimensional action

$$p(x) = \frac{1}{(2\pi)^{k/2} \prod_{i=1}^k \sigma_i} \exp\left(-\frac{1}{2} \sum_{i=1}^k \frac{(x_i - \mu_i)^2}{\sigma_i^2}\right) \quad \log p(x) = -\frac{1}{2} \sum_{i=1}^k \left[ \frac{(x_i - \mu_i)^2}{\sigma_i^2} + \log(2\pi\sigma_i^2) \right]$$

- Log-likelihood

$$\log \pi_\theta(a|s) = -\frac{1}{2} \left( \sum_{i=1}^k \left( \frac{(a_i - \mu_i)^2}{\sigma_i^2} + 2 \log \sigma_i \right) + k \log 2\pi \right)$$

# RL Concepts

- A Trajectory is a sequence of states and actions in the world

$$\tau = (s_0, a_0, s_1, a_1, \dots)$$

- Start-state distribution  $s_0 \sim \rho_0(\cdot)$
- State transitions are governed by natural laws of the environment (dynamics)

- Deterministic  $s_{t+1} = f(s_t, a_t)$

- Stochastic  $s_{t+1} \sim P(\cdot | s_t, a_t)$

# RL Concepts

simplified

- Reward function  $r_t = R(s_t, a_t, s_{t+1})$   $r_t = R(s_t)$   $r_t = R(s_t, a_t)$

- Finite-horizon undiscounted return for a trajectory

$$R(\tau) = \sum_{t=0}^T r_t$$

- Infinite-horizon discounted return for a trajectory

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t \quad \gamma \in (0, 1)$$

# The RL Problem

- The goal of RL is to select a policy which **maximizes expected return** when the agent acts according to it
- Probability distribution over trajectories

$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$$

- Expected return  $J(\pi) = \int_{\tau} P(\tau|\pi) R(\tau) = \mathbb{E}_{\tau \sim \pi} [R(\tau)]$  Sample trajectories

- The central optimization problem  $\pi^* = \arg \max_{\pi} J(\pi)$  Optimal policy

In practice

$$\pi_{\theta}^* = \arg \max_{\theta} J(\pi_{\theta})$$

Learn the parameters  
of the policy

# Summary

- RL concepts

# Further Reading

- OpenAI Spinning Up in Deep RL

<https://spinningup.openai.com/en/latest/index.html>