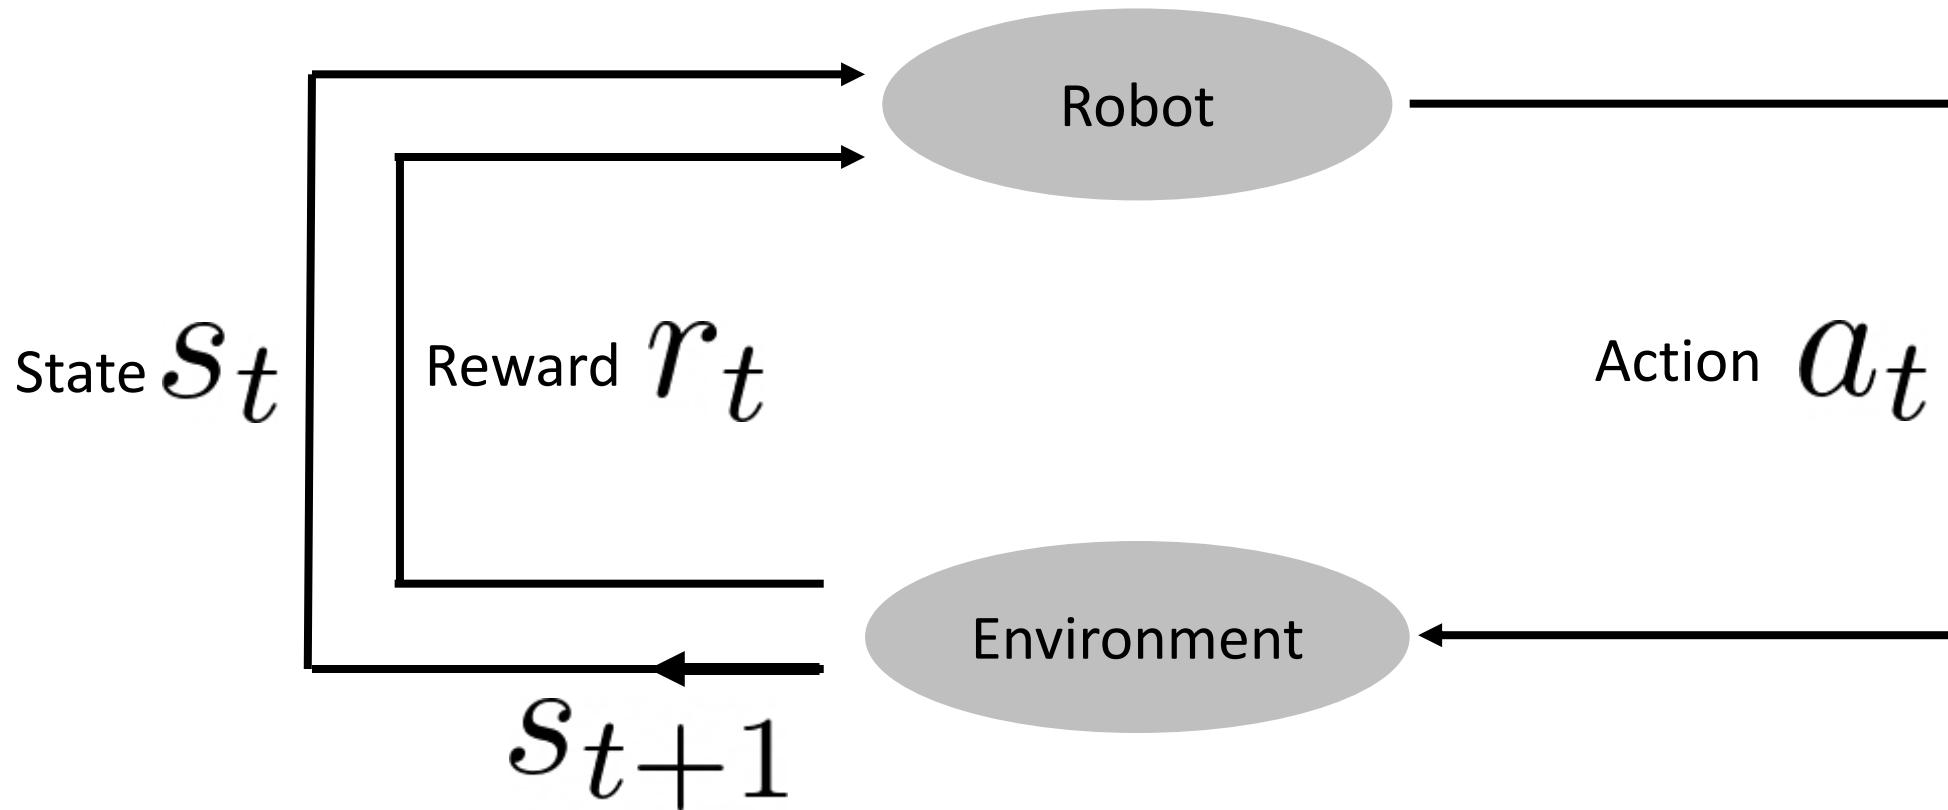




Reinforcement Learning: Policy Optimization

CS 6341 Robotics
Professor Yu Xiang
The University of Texas at Dallas

Reinforcement Learning



Reinforcement Learning:
Imitation Learning:

$$a_t = \pi(s_t)$$

The RL Problem

- The goal of RL is to select a policy which **maximizes expected return** when the agent acts according to it
- Probability distribution over trajectories

$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$$

Transition model (no
need in model-free RL)

$$p(s' | s, a)$$

Sample trajectories

$$R(\tau) = \sum_{t=0}^T r_t \quad R(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t$$
$$\gamma \in (0, 1)$$

- Expected return

$$J(\pi) = \int_{\tau} P(\tau|\pi) R(\tau) = \mathbb{E}_{\tau \sim \pi} [R(\tau)]$$

- The central optimization problem

$$\pi^* = \arg \max_{\pi} J(\pi)$$

Optimal policy

In practice

$$\pi_{\theta}^* = \arg \max_{\theta} J(\pi_{\theta})$$

Learn the parameters
of the policy

Value Functions

- Value of a state or a state-action pair
 - The expected return if you start in that state or state-action pair, and then act according to a particular policy forever after
- On-policy Value Function $V^\pi(s) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s]$
- On-policy Action-Value Function $Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s, a_0 = a]$
- Optimal Value Function $V^*(s) = \max_{\pi} \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s]$
- Optimal Action-Value Function $Q^*(s, a) = \max_{\pi} \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s, a_0 = a]$

Value Functions

- Connection

$$V^\pi(s) = \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)]$$

$$V^*(s) = \max_a Q^*(s, a)$$

- The optimal policy in \mathcal{S} will select whichever action maximizes the expected return starting in \mathcal{S}

$$a^*(s) = \arg \max_a Q^*(s, a)$$

Parametrized Value Functions

- On-policy Value Function $V^\pi(s) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s]$
- Parameterization (a network) $V_\phi(s)$
- Learning the value function
 - Sample trajectories
 - For each trajectory $G_t = \sum_{k=t}^T \gamma^{k-t} r_k$
 - Supervised learning $L(\phi) = \frac{1}{N} \sum_t (V_\phi(s_t) - G_t)^2$

Bellman Equations

- The value of your starting point is the reward you expect to get from being there, plus the value of wherever you land next

$$V^\pi(s) = \underset{\substack{a \sim \pi \\ s' \sim P}}{\text{E}} [r(s, a) + \gamma V^\pi(s')],$$

On-policy

$$Q^\pi(s, a) = \underset{s' \sim P}{\text{E}} \left[r(s, a) + \gamma \underset{a' \sim \pi}{\text{E}} [Q^\pi(s', a')] \right]$$

$$V^*(s) = \max_a \underset{s' \sim P}{\text{E}} [r(s, a) + \gamma V^*(s')],$$

Optimal policy

$$Q^*(s, a) = \underset{s' \sim P}{\text{E}} \left[r(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

Advantage Functions

- How much better it is to take a specific action a in state s , over randomly selecting an action according to $\pi(\cdot|s)$

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

Q value for (s, a)

V value for s : random action from the policy

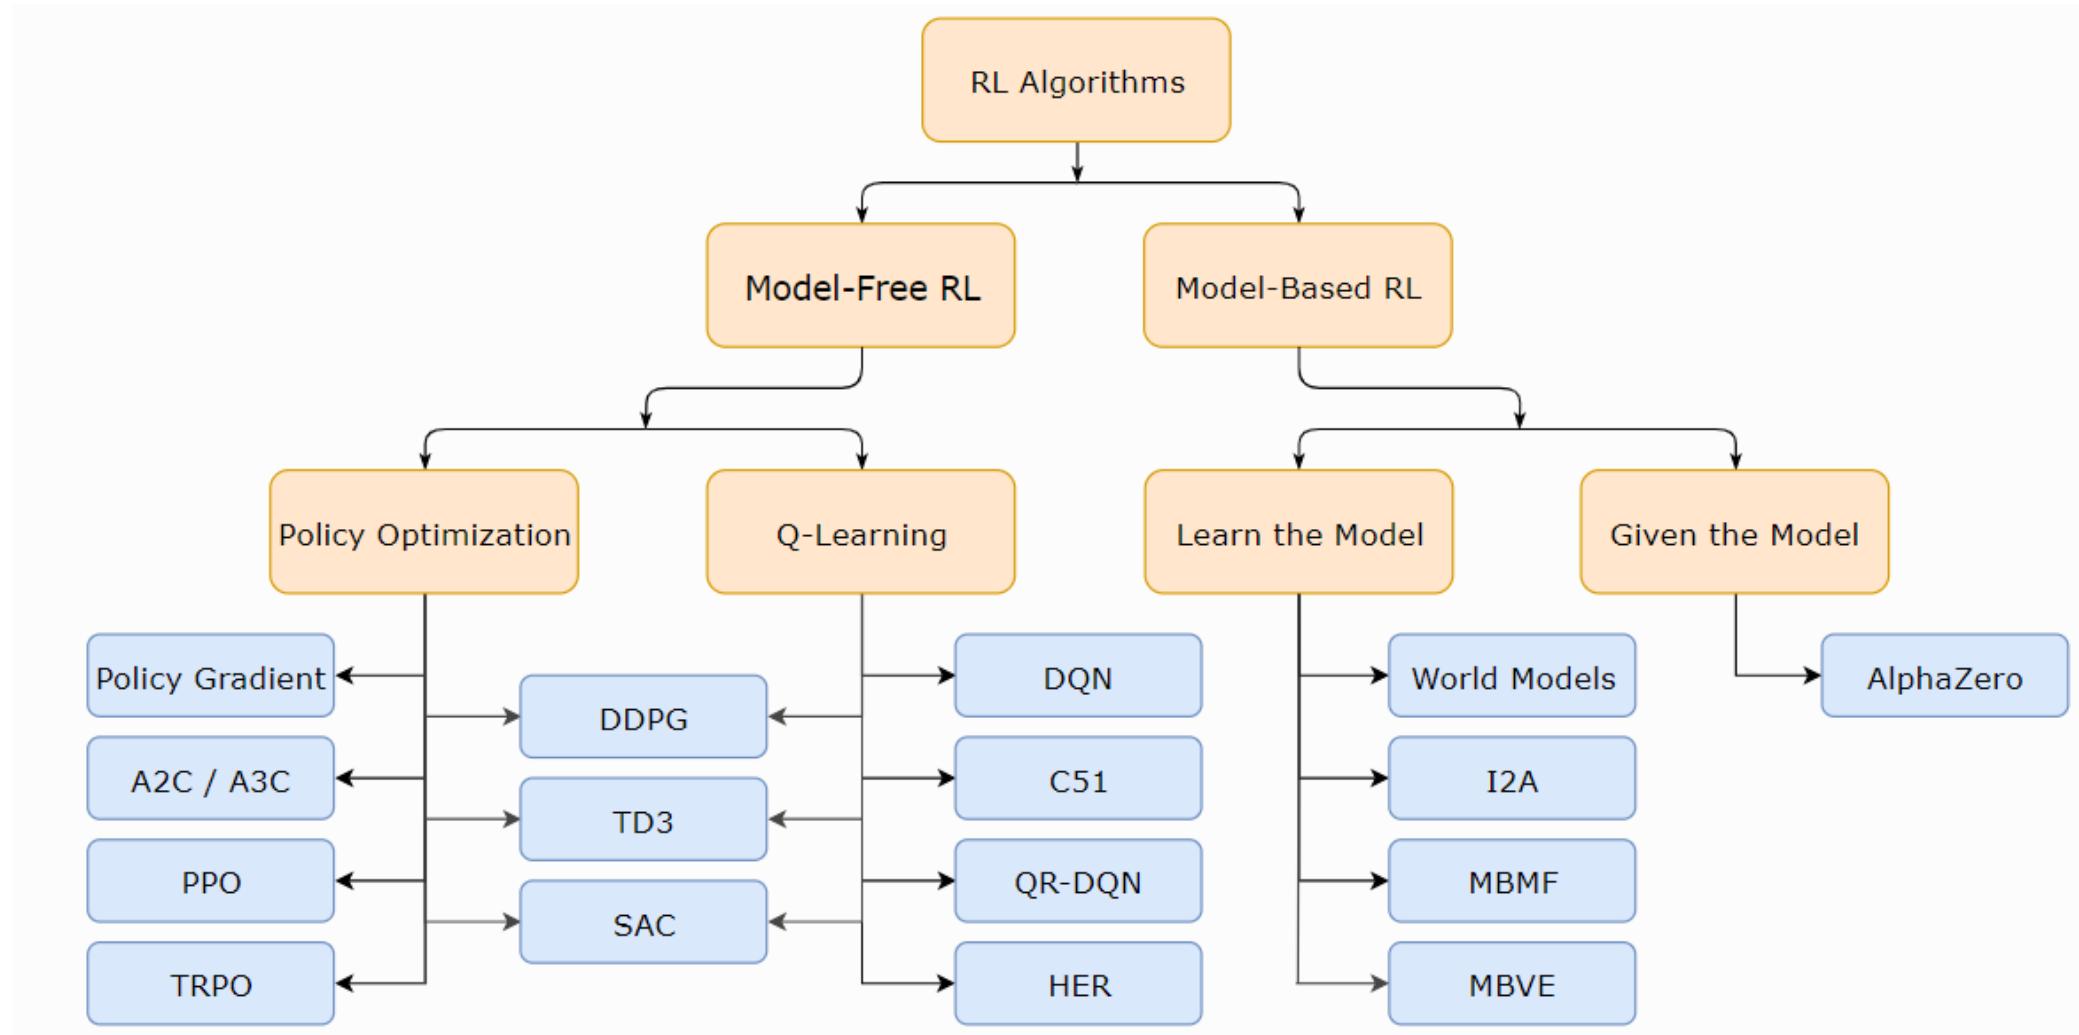
$$Q^\pi(s, a) = \underset{\tau \sim \pi}{\text{E}} [R(\tau) | s_0 = s, a_0 = a]$$

$$V^\pi(s) = \underset{\tau \sim \pi}{\text{E}} [R(\tau) | s_0 = s]$$

Markov Decision Processes (MDPs)

- A MDP is a 5-tuple $\langle S, A, R, P, \rho_0 \rangle$
 - S is the set of all valid states,
 - A is the set of all valid actions,
 - $R : S \times A \times S \rightarrow \mathbb{R}$ is the reward function, with $r_t = R(s_t, a_t, s_{t+1})$,
 - $P : S \times A \rightarrow \mathcal{P}(S)$ is the transition probability function, with $P(s'|s, a)$ being the probability of transitioning into state s' if you start in state s and take action a ,
 - and ρ_0 is the starting state distribution.

A Taxonomy of RL Algorithms



Model-Free vs. Model-based RL

- Whether the agent has access to (or learns) a model of the environment
- A model is a function which predicts state transitions and reward
 - Transition model $p(s' | s, a)$
 - Reward model $r(s, a)$
- A model allows the agent to plan by thinking ahead
- A ground-truth model of the environment is usually not available to the agent

Model-Free RL: Policy Gradient

- Maximize expected return
- Gradient ascent

$$J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$$

$$R(\tau) = \sum_{t=0}^T r_t$$

- How to compute the policy gradient?

$$\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_\theta)|_{\theta_k}$$

Policy gradient

$$\begin{aligned}\nabla_\theta J(\pi_\theta) &= \nabla_\theta \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)] \\ &= \nabla_\theta \int_{\tau} P(\tau|\theta) R(\tau) \\ &= \int_{\tau} \nabla_\theta P(\tau|\theta) R(\tau)\end{aligned}$$

Probability of a Trajectory

$$P(\tau|\theta) = \rho_0(s_0) \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

Policy Gradient

- The Log-Derivative Trick

$$\nabla_{\theta} P(\tau|\theta) = P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta)$$

$$\log P(\tau|\theta) = \log \rho_0(s_0) + \sum_{t=0}^T \left(\log P(s_{t+1}|s_t, a_t) + \log \pi_{\theta}(a_t|s_t) \right)$$

$$\nabla_{\theta} \log P(\tau|\theta) = \cancel{\nabla_{\theta} \log \rho_0(s_0)} + \sum_{t=0}^T \left(\cancel{\nabla_{\theta} \log P(s_{t+1}|s_t, a_t)} + \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \right)$$

$$= \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t|s_t).$$

No need to know the transition model

Policy Gradient

$$\begin{aligned}\nabla_{\theta} J(\pi_{\theta}) &= \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] \\&= \nabla_{\theta} \int_{\tau} P(\tau | \theta) R(\tau) && \text{Expand expectation} \\&= \int_{\tau} \nabla_{\theta} P(\tau | \theta) R(\tau) && \text{Bring gradient under integral} \\&= \int_{\tau} P(\tau | \theta) \nabla_{\theta} \log P(\tau | \theta) R(\tau) && \text{Log-derivative trick} \\&= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau | \theta) R(\tau)] && \text{Return to expectation form}\end{aligned}$$

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R(\tau) \right] \quad \text{Expression for grad-log-prob}$$

Policy Gradient

- Collect a set of trajectories using the policy π_θ

$$\mathcal{D} = \{\tau_i\}_{i=1,\dots,N}$$

- Estimate policy gradient

$$\hat{g} = \frac{1}{|\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) R(\tau)$$

Categorical policy for discrete actions

$$\log \pi_\theta(a|s) = \log [P_\theta(s)]_a$$

Diagonal Gaussian policy

$$\log \pi_\theta(a|s) = -\frac{1}{2} \left(\sum_{i=1}^k \left(\frac{(a_i - \mu_i)^2}{\sigma_i^2} + 2 \log \sigma_i \right) + k \log 2\pi \right)$$

Policy Gradient

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R(\tau) \right]$$

$$R(\tau) = \sum_{t=0}^T r_t$$

Agents should really only reinforce actions on the basis of their consequences.

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \sum_{t'=t}^T R(s_{t'}, a_{t'}, s_{t'+1}) \right]$$

$$\hat{R}_t \doteq \sum_{t'=t}^T R(s_{t'}, a_{t'}, s_{t'+1}) \quad \text{reward-to-go}$$

Vanilla Policy Gradient

- Key idea: push up the probabilities of actions that lead to higher return, and push down probabilities of actions that lead to lower return
- The expected finite-horizon undiscounted return of the policy $J(\pi_\theta)$

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t | s_t) A^{\pi_\theta}(s_t, a_t) \right]$$

Advantage function $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$

Stochastic gradient ascent $\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_{\theta_k})$

Vanilla Policy Gradient

Algorithm 1 Vanilla Policy Gradient Algorithm

- 1: Input: initial policy parameters θ_0 , initial value function parameters ϕ_0
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_k = \pi(\theta_k)$ in the environment.
- 4: Compute rewards-to-go \hat{R}_t .
- 5: Compute advantage estimates, \hat{A}_t (using any method of advantage estimation) based on the current value function V_{ϕ_k} .
- 6: Estimate policy gradient as

$$\hat{g}_k = \frac{1}{|\mathcal{D}_k|} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) |_{\theta_k} \hat{A}_t.$$

- 7: Compute policy update, either using standard gradient ascent,

$$\theta_{k+1} = \theta_k + \alpha_k \hat{g}_k,$$

or via another gradient ascent algorithm like Adam.

- 8: Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k| T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 9: **end for**
-

Exploration vs. Exploitation

- stochastic policy

reward-to-go

$$\hat{R}_t \doteq \sum_{t'=t}^T R(s_{t'}, a_{t'}, s_{t'+1})$$

Advantage function

$$\begin{aligned} A^{\pi}(s, a) &= Q^{\pi}(s, a) - V^{\pi}(s) \\ &= r + V^{\pi}(s') - V^{\pi}(s) \end{aligned}$$

Bellman Equations

Trust Region Policy Optimization (TRPO)

- TRPO update

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} \mathcal{L}(\theta_k, \theta) \\ \text{s.t. } \bar{D}_{KL}(\theta || \theta_k) &\leq \delta\end{aligned}$$

taking the largest step possible
to improve performance

- *Surrogate advantage*

$$\mathcal{L}(\theta_k, \theta) = \underset{s, a \sim \pi_{\theta_k}}{\mathbb{E}} \left[\frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right]$$

A measure of how the policy performs related to the old policy

- *KL-divergence*

$$\bar{D}_{KL}(\theta || \theta_k) = \underset{s \sim \pi_{\theta_k}}{\mathbb{E}} [D_{KL} (\pi_\theta(\cdot|s) || \pi_{\theta_k}(\cdot|s))]$$

Proximal Policy Optimization (PPO)

- PPO-clip updates $\theta_{k+1} = \arg \max_{\theta} \mathbb{E}_{s,a \sim \pi_{\theta_k}} [L(s, a, \theta_k, \theta)]$

$$L(s, a, \theta_k, \theta) = \min \left(\frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a), \text{clip} \left(\frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)}, 1 - \epsilon, 1 + \epsilon \right) A^{\pi_{\theta_k}}(s, a) \right)$$

Avoid stepping so far that we accidentally cause performance collapse

PPO methods are significantly simpler to implement, and empirically seem to perform at least as well as TRPO

- A simpler version

$$L(s, a, \theta_k, \theta) = \min \left(\frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a), g(\epsilon, A^{\pi_{\theta_k}}(s, a)) \right)$$

$$g(\epsilon, A) = \begin{cases} (1 + \epsilon)A & A \geq 0 \\ (1 - \epsilon)A & A < 0. \end{cases}$$

Proximal Policy Optimization (PPO)

Algorithm 1 PPO-Clip

- 1: Input: initial policy parameters θ_0 , initial value function parameters ϕ_0
- 2: **for** $k = 0, 1, 2, \dots$ **do**
- 3: Collect set of trajectories $\mathcal{D}_k = \{\tau_i\}$ by running policy $\pi_k = \pi(\theta_k)$ in the environment.
- 4: Compute rewards-to-go \hat{R}_t .
- 5: Compute advantage estimates, \hat{A}_t (using any method of advantage estimation) based on the current value function V_{ϕ_k} .
- 6: Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left(\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_k}(a_t | s_t)} A^{\pi_{\theta_k}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

- 7: Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left(V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 8: **end for**
-

PPO Example



<https://rewind-reward.github.io/>

Summary

- Model-free RL
 - Vanilla Policy Gradient
 - Trust Region Policy Optimization (TRPO)
 - Proximal Policy Optimization (PPO)

Further Reading

- OpenAI Spinning Up in Deep RL

<https://spinningup.openai.com/en/latest/index.html>