



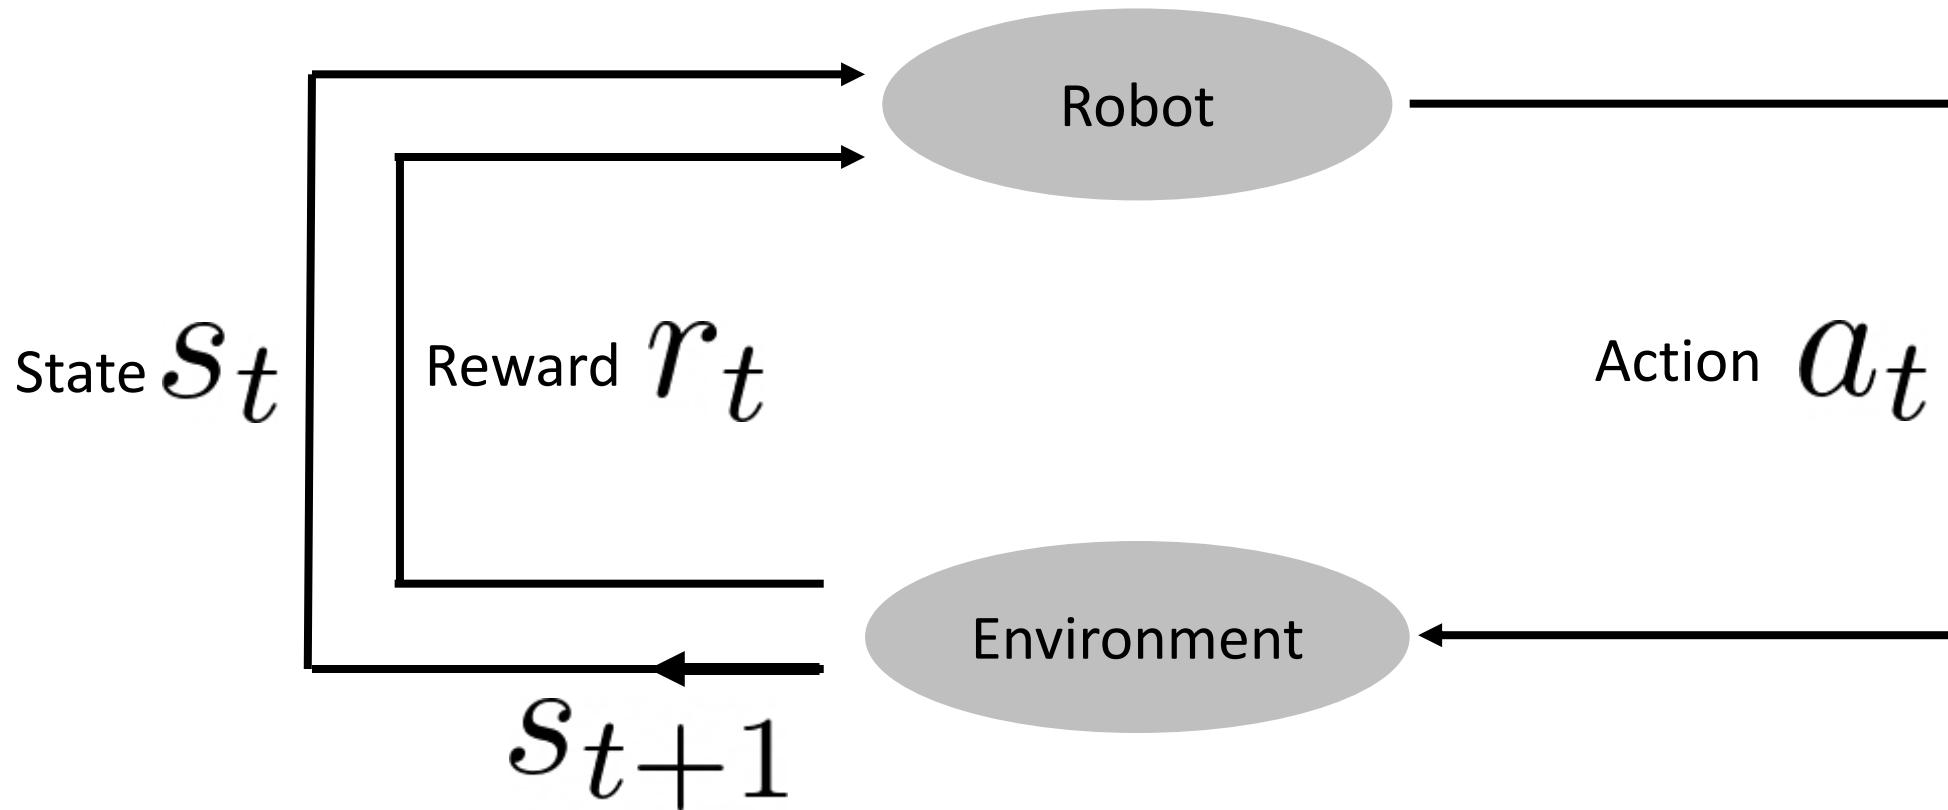
# Reinforcement Learning: Actor-Critic

CS 6341 Robotics

Professor Yu Xiang

The University of Texas at Dallas

# Reinforcement Learning



Reinforcement Learning:  
Imitation Learning:

$$a_t = \pi(s_t)$$

# Last Lecture: Policy Optimization

- Maximize expected return

$$J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$$

$$R(\tau) = \sum_{t=0}^T r_t$$

$$J(\pi_\theta) = \int_{\tau} P(\tau|\theta) R(\tau)$$

$$P(\tau|\theta) = \rho_0(s_0) \prod_{t=0}^T P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)$$

- Policy gradient

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) R(\tau) \right]$$

$$\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_\theta)|_{\theta_k}$$

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) \sum_{t'=t}^T R(s_{t'}, a_{t'}, s_{t'+1}) \right]$$

reward-to-go

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) A^{\pi_\theta}(s_t, a_t) \right]$$

Advantage

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

# Q-Learning

- Learn the optimal Q function

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s, a_0 = a]$$

- Policy from the Q function

$$a^*(s) = \arg \max_a Q^*(s, a)$$

- How to learn the Q function?

- Bellman Equation

$$Q^*(s, a) = \mathbb{E}_{s' \sim P} [r(s, a) + \gamma \max_{a'} Q^*(s', a')]$$

# Q-Learning

- For discrete states and actions
- Dynamic programming (Q-table)
  - Initialize Q values arbitrarily  $Q_0(s, a) = 0$
  - Then iterate

$$Q_{k+1}(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a')$$

$$a^*(s) = \arg \max_a Q^*(s, a)$$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

<https://mohitmayank.com/blog/interactive-q-learning>

# Q-Learning

$$Q^*(s, a) = \underset{s' \sim P}{\mathbb{E}} \left[ r(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

- What if the state and action space is large?
  - We cannot store a table
- Use parameterization  $Q_\phi(s, a)$
- Collect a set of transitions  $\mathcal{D} = \{(s_i, a_i, r_i, s'_i)\}$  Replay Buffer
- TD target  $y_i = r_i + \gamma \max_{a'} Q_{\phi^-}(s'_i, a')$
- Loss function  $L(\phi) = \frac{1}{N} \sum_i (Q_\phi(s_i, a_i) - y_i)^2$   $\phi \leftarrow \phi - \alpha \nabla_\phi L(\phi)$
- Update the target network  $\phi^- \leftarrow \phi$

# Q-Learning

- TD target  $y_i = r_i + \gamma \max_{a'} Q_{\phi^-}(s'_i, a')$
- How to compute this max?
  - Discretize actions  $Q_{\phi^-}(s') = [Q_{\phi^-}(s', a_1), Q_{\phi^-}(s', a_2), \dots, Q_{\phi^-}(s', a_K)]$

---

## Playing Atari with Deep Reinforcement Learning

---

Volodymyr Mnih   Koray Kavukcuoglu   David Silver   Alex Graves   Ioannis Antonoglou

Daan Wierstra   Martin Riedmiller

DeepMind Technologies

We now describe the exact architecture used for all seven Atari games. The input to the neural network consists is an  $84 \times 84 \times 4$  image produced by  $\phi$ . The first hidden layer convolves 16  $8 \times 8$  filters with stride 4 with the input image and applies a rectifier nonlinearity [10, 18]. The second hidden layer convolves 32  $4 \times 4$  filters with stride 2, again followed by a rectifier nonlinearity. The final hidden layer is fully-connected and consists of 256 rectifier units. The output layer is a fully-connected linear layer with a single output for each valid action. The number of valid actions varied between 4 and 18 on the games we considered. We refer to convolutional networks trained with our approach as Deep Q-Networks (DQN).

Volodymyr Mnih et al., 2013 (arXiv preprint)

# Q-Learning

- TD target  $y_i = r_i + \gamma \max_{a'} Q_{\phi^-}(s'_i, a')$
- How to compute this max?
  - Discretize actions  $Q_{\phi^-}(s') = [Q_{\phi^-}(s', a_1), Q_{\phi^-}(s', a_2), \dots, Q_{\phi^-}(s', a_K)]$
  - Continuous actions: actor-critic methods
    - Learn a policy (actor)  $\pi_\theta(s') \approx \arg \max_a Q_\phi(s', a)$
    - $$y = r + \gamma Q_{\phi^-}(s', \pi_\theta(s'))$$

# Deep Deterministic Policy Gradient (DDPG)

- DDPG currently learns a Q-function and a policy
  - Uses off-policy data and the Bellman equation to learn the Q-function
  - Uses the Q-function to learn the policy

- Q-learning

$$Q^*(s, a) = \underset{s' \sim P}{\text{E}} \left[ r(s, a) + \gamma \max_{a'} Q^*(s', a') \right]$$

Approximator  $Q_\phi(s, a)$       Collect a set of transitions  $(s, a, r, s', d)$

**mean-squared  
Bellman error  
(MSBE)**

$$L(\phi, \mathcal{D}) = \underset{(s, a, r, s', d) \sim \mathcal{D}}{\text{E}} \left[ \left( Q_\phi(s, a) - \left( r + \gamma(1 - d) \max_{a'} Q_\phi(s', a') \right) \right)^2 \right]$$

a policy  $\mu(s)$

$\max_a Q(s, a) \approx Q(s, \mu(s))$

$Q_\phi(s', \mu(s'))$

# Deep Deterministic Policy Gradient (DDPG)

- Trick one: replay buffers
  - Large enough to contain a wide range of experiences
- Trick two: target networks
  - The term is called target  $r + \gamma(1 - d) \max_{a'} Q_\phi(s', a')$
  - The target depends on the same parameters  $\phi$ , but with a time delay
  - Target network  $\phi_{\text{targ}}$ 
$$\phi_{\text{targ}} \leftarrow \rho\phi_{\text{targ}} + (1 - \rho)\phi$$
  - Target policy network  $\mu_{\theta_{\text{targ}}}$

# Deep Deterministic Policy Gradient (DDPG)

- Q-learning in DDPG

$$L(\phi, \mathcal{D}) = \underset{(s,a,r,s',d) \sim \mathcal{D}}{\text{E}} \left[ \left( Q_\phi(s, a) - (r + \gamma(1-d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))) \right)^2 \right]$$

- Policy learning in DDPG

$$\max_{\theta} \underset{s \sim \mathcal{D}}{\text{E}} [Q_\phi(s, \mu_\theta(s))]$$

Gradient Ascent

# Deep Deterministic Policy Gradient (DDPG)

---

**Algorithm 1** Deep Deterministic Policy Gradient

---

```
1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi$ , empty replay buffer  $\mathcal{D}$ 
2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta$ ,  $\phi_{\text{targ}} \leftarrow \phi$ 
3: repeat
4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_\theta(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$ 
5:   Execute  $a$  in the environment
6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$ 
8:   If  $s'$  is terminal, reset environment state.
9:   if it's time to update then
10:    for however many updates do
11:      Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$ 
12:      Compute targets
13:      Update Q-function by one step of gradient descent using
14:        
$$\nabla_\phi \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_\phi(s, a) - y(r, s', d))^2$$

15:        Update policy by one step of gradient ascent using
16:          
$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} Q_\phi(s, \mu_\theta(s))$$

17:        Update target networks with
18:          
$$\phi_{\text{targ}} \leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi$$

          
$$\theta_{\text{targ}} \leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta$$

16:      end for
17:      end if
18:    until convergence
```

---

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

13: Update Q-function by one step of gradient descent using

$$\nabla_\phi \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_\phi(s, a) - y(r, s', d))^2$$

# Twin Delayed DDPG (TD3)

- Trick one: clipped double-Q learning
  - TD3 learns two Q functions Q-learning suffers from overestimation bias
  - uses the smaller of the two Q-values to form the targets in the Bellman error loss functions
- Trick two: “delayed” policy updates
  - Updates the policy (and target networks) less frequently than the Q-function
- Trick three: target policy smoothing
  - Adds noise to the target action, to make it harder for the policy to exploit Q-function errors by smoothing out Q along changes in action

$$a'(s') = \text{clip} \left( \mu_{\theta_{\text{targ}}}(s') + \text{clip}(\epsilon, -c, c), a_{\text{Low}}, a_{\text{High}} \right), \quad \epsilon \sim \mathcal{N}(0, \sigma)$$

# Soft Actor-Critic (SAC)

$$\log \pi_\theta(a|s) = \log [P_\theta(s)]_a$$

$$\log \pi_\theta(a|s) = -\frac{1}{2} \left( \sum_{i=1}^k \left( \frac{(a_i - \mu_i)^2}{\sigma_i^2} + 2 \log \sigma_i \right) + k \log 2\pi \right)$$

- An algorithm that optimizes a stochastic policy in an off-policy way
- Entropy-regularized RL

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t \left( R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot|s_t)) \right) \right]$$

$$\text{Entropy } H(P) = \mathbb{E}_{x \sim P} [-\log P(x)]$$

increasing entropy results in more exploration, which can accelerate learning later on

$$V^\pi(s) = \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t \left( R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot|s_t)) \right) \middle| s_0 = s \right]$$

$$V^\pi(s) = \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s))$$

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) + \alpha \sum_{t=1}^{\infty} \gamma^t H(\pi(\cdot|s_t)) \middle| s_0 = s, a_0 = a \right]$$

# Soft Actor-Critic (SAC)

- SAC learns a policy and two Q-functions
  - Uses entropy regularization
  - Train a stochastic policy

$$\begin{aligned}Q^\pi(s, a) &= \underset{\substack{s' \sim P \\ a' \sim \pi}}{\text{E}} [R(s, a, s') + \gamma (Q^\pi(s', a') + \alpha H(\pi(\cdot|s')))] \\&= \underset{\substack{s' \sim P \\ a' \sim \pi}}{\text{E}} [R(s, a, s') + \gamma (Q^\pi(s', a') - \alpha \log \pi(a'|s'))]\end{aligned}$$

Approximate expectation with samples  $Q^\pi(s, a) \approx r + \gamma (Q^\pi(s', \tilde{a}') - \alpha \log \pi(\tilde{a}'|s')) , \quad \tilde{a}' \sim \pi(\cdot|s')$

# Soft Actor-Critic (SAC)

- Q-learning

$$L(\phi_i, \mathcal{D}) = \underset{(s,a,r,s',d) \sim \mathcal{D}}{\text{E}} \left[ \left( Q_{\phi_i}(s, a) - y(r, s', d) \right)^2 \right]$$

$$y(r, s', d) = r + \gamma(1 - d) \left( \min_{j=1,2} Q_{\phi_{\text{targ},j}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

- Policy learning

The policy is learned by maximizing the **soft value function**

$$\begin{aligned} \text{maximize } V^\pi(s) &= \underset{a \sim \pi}{\text{E}} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s)) \\ &= \underset{a \sim \pi}{\text{E}} [Q^\pi(s, a) - \alpha \log \pi(a|s)] \end{aligned}$$

# Soft Actor-Critic (SAC)

- Policy learning

**reparameterization trick**       $\tilde{a}_\theta(s, \xi) = \tanh(\mu_\theta(s) + \sigma_\theta(s) \odot \xi), \quad \xi \sim \mathcal{N}(0, I)$

$$\mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(s, a) - \alpha \log \pi_\theta(a|s)] = \mathbb{E}_{\xi \sim \mathcal{N}} [Q^{\pi_\theta}(s, \tilde{a}_\theta(s, \xi)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s, \xi)|s)]$$

$$\max_{\theta} \mathbb{E}_{\substack{s \sim \mathcal{D} \\ \xi \sim \mathcal{N}}} \left[ \min_{j=1,2} Q_{\phi_j}(s, \tilde{a}_\theta(s, \xi)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s, \xi)|s) \right]$$

# Soft Actor-Critic (SAC)

---

**Algorithm 1** Soft Actor-Critic

```

1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi_1, \phi_2$ , empty replay buffer  $\mathcal{D}$ 
2: Set target parameters equal to main parameters  $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$ 
3: repeat
4:   Observe state  $s$  and select action  $a \sim \pi_\theta(\cdot|s)$ 
5:   Execute  $a$  in the environment
6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$ 
8:   If  $s'$  is terminal, reset environment state.
9:   if it's time to update then
10:    for  $j$  in range(however many updates) do
11:      Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$ 
12:      Compute targets for the Q functions:
```

$$y(r, s', d) = r + \gamma(1 - d) \left( \min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

---

```

13:   Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

14:   Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left( \min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$


where  $\tilde{a}_\theta(s)$  is a sample from  $\pi_\theta(\cdot|s)$  which is differentiable wrt  $\theta$  via the reparametrization trick.
15:   Update target networks with

$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

16:   end for
17:   end if
18: until convergence
```


---

# Learning Closed-Loop Control Policies for 6D Grasping



Goal-Auxiliary Actor-Critic for 6D Robotic Grasping with Point Clouds. Wang-Xiang-Yang-Mousavian-Fox, CoRL'21






# GA-DDPG Network Architecture





# Learning from Demonstration with the OMG-

50,000 trajectories  
planar


1,500 3D shapes



# Our Learned Policy in the Real World



# Closed-Loop Human-Robot Handover



# Summary

- Model-free RL
  - Deep Deterministic Policy Gradient (DDPG)
  - Twin Delayed DDPG (TD3)
  - Soft Actor-Critic (SAC)

# Further Reading

- OpenAI Spinning Up in Deep RL

<https://spinningup.openai.com/en/latest/index.html>